Computer Networks
CS3611

Introduction-Part 2

Haiming Jin

The slides are adapted from those provided by Prof. Shizhen Zhao and Romit Roy Choudhury.

Chapter 1: Introduction

Our goal:

- get "feel" and terminology
- paint a broad picture
- see the forest through the trees
- approach:
  - use Internet as example

Topics:

- What's Computer Network?
- Protocol layers, service models
- What's the Internet?
- Network edge
- Access net and physical media
- Network core
- Internet structure and ISPs
- Delay, loss, and throughput in packet-switched networks
- History of Internet

Chapter 1: roadmap

- What's Computer Network?
- Protocol layers, service models
- What's the Internet?
- Network edge
- Access net and physical media
- Network core
- Internet structure and ISPs
- Delay, loss, and throughput in packet-switched networks
- History of Internet

What's the Internet: "nuts and bolts" view

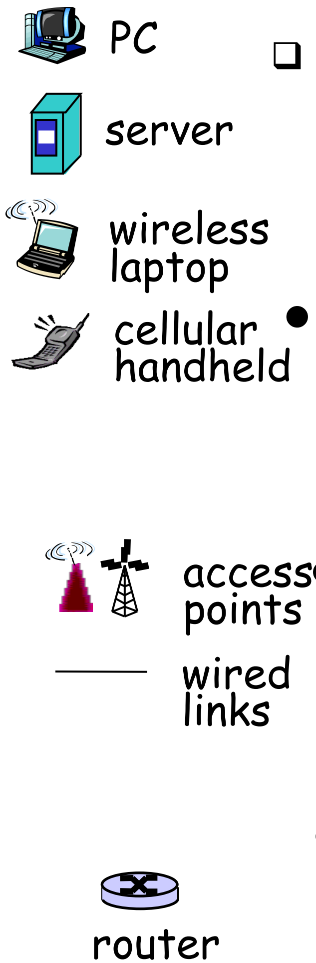

- millions of connected computing devices: *hosts , end systems*
- *communication links*
  - fiber, copper, radio, satellite
- *routers:* forward packets (chunks of data)
- *protocols* control sending, receiving of msgs
  - e.g., TCP, IP, HTTP

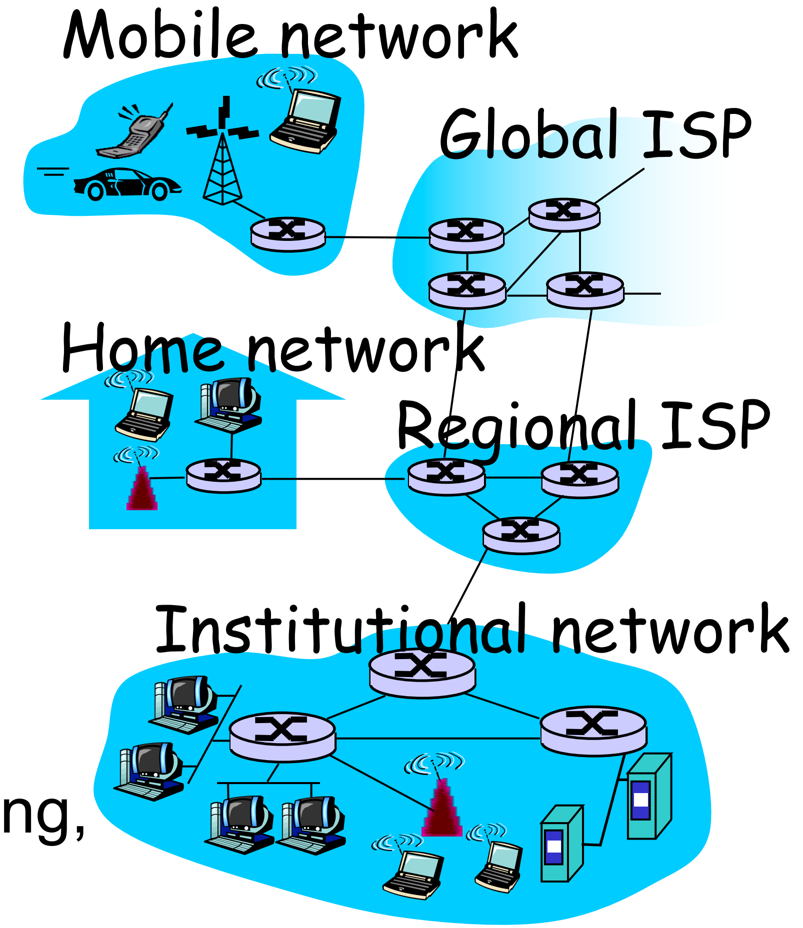

What’s the Internet: a service view

- communication *infrastructure* enables distributed applications:
  - Web, VoIP, email, games, e-commerce, file sharing
- communication services provided to apps:
  - reliable data delivery from source to destination
  - “best effort” (unreliable) data delivery

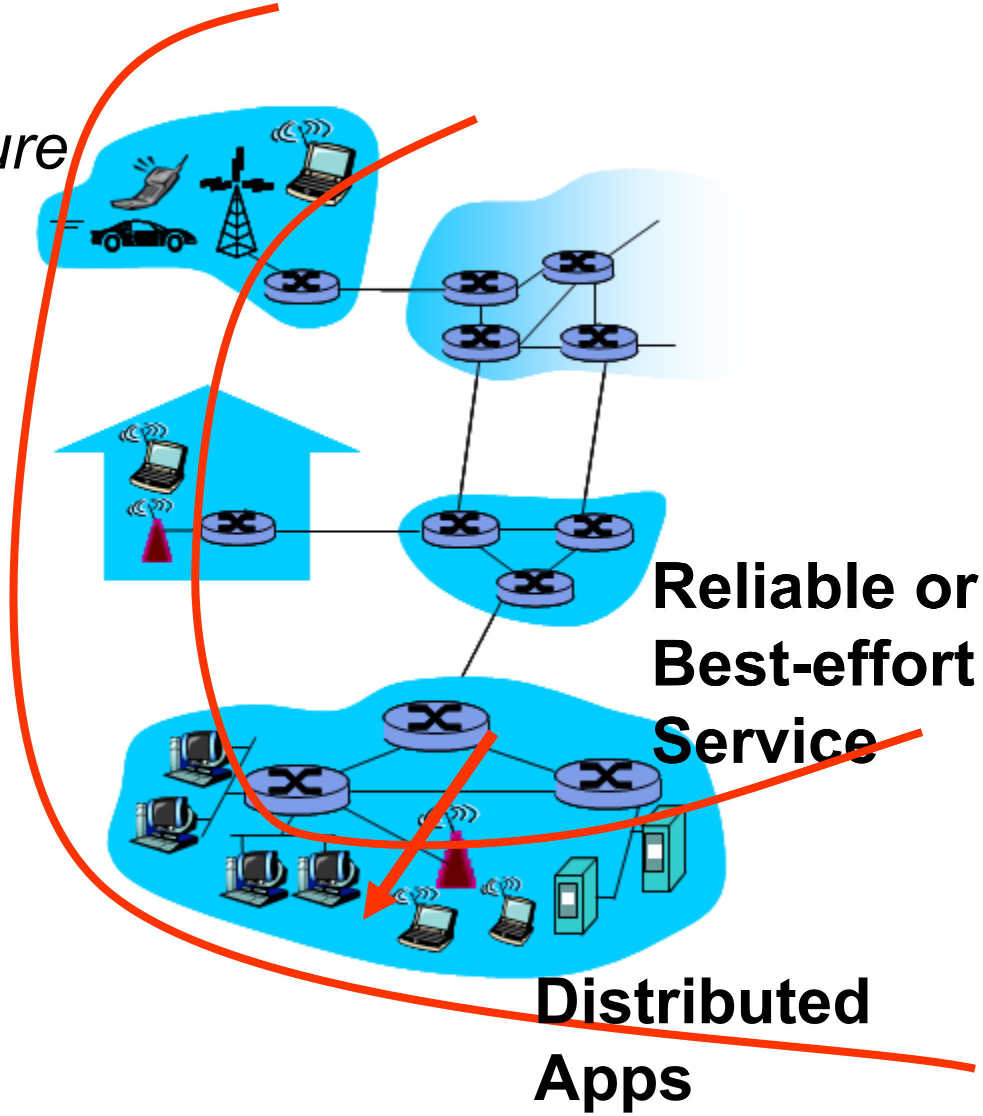

Internet Structure

- network edge: applications and hosts
- access networks, physical media: wired, wireless communication links
- network core:
  - interconnected routers
  - network of networks

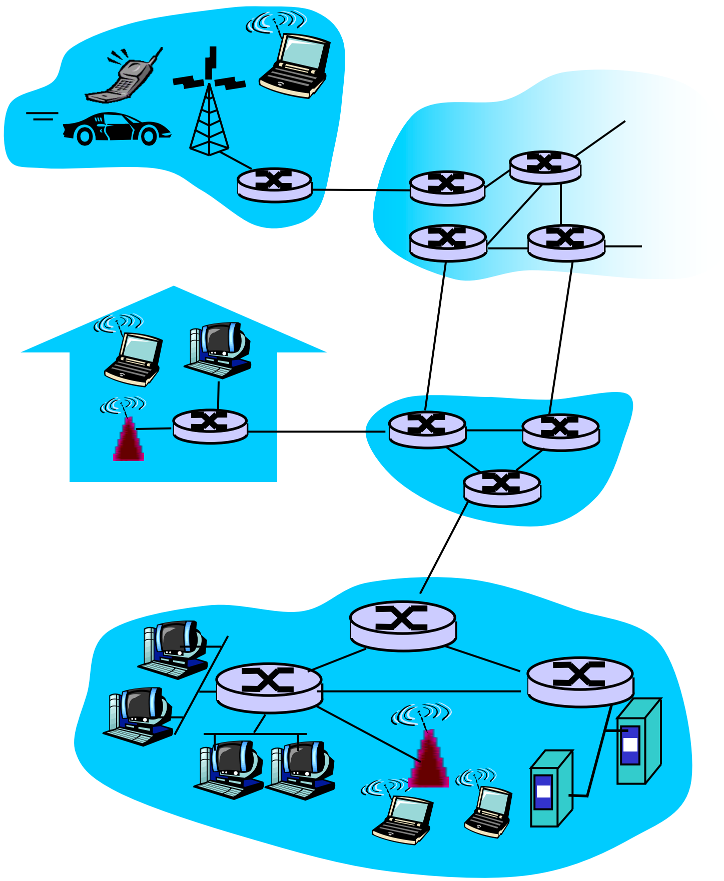

Chapter 1: roadmap

- What's Computer Network?
- Protocol layers, service models
- What's the Internet?
- Network edge
- Access net and physical media
- Network core
- Internet structure and ISPs
- Delay, loss, and throughput in packet-switched networks
- History of Internet

The network edge

- end systems (hosts):
  - run application programs
  - e.g. Web, email
  - at "edge of network"
- client/server model
  - client host requests, receives service from always-on server
  - e.g. Web browser/server; email client/server
- peer-peer model:
  - minimal (or no) use of dedicated servers
  - e.g. Skype, BitTorrent

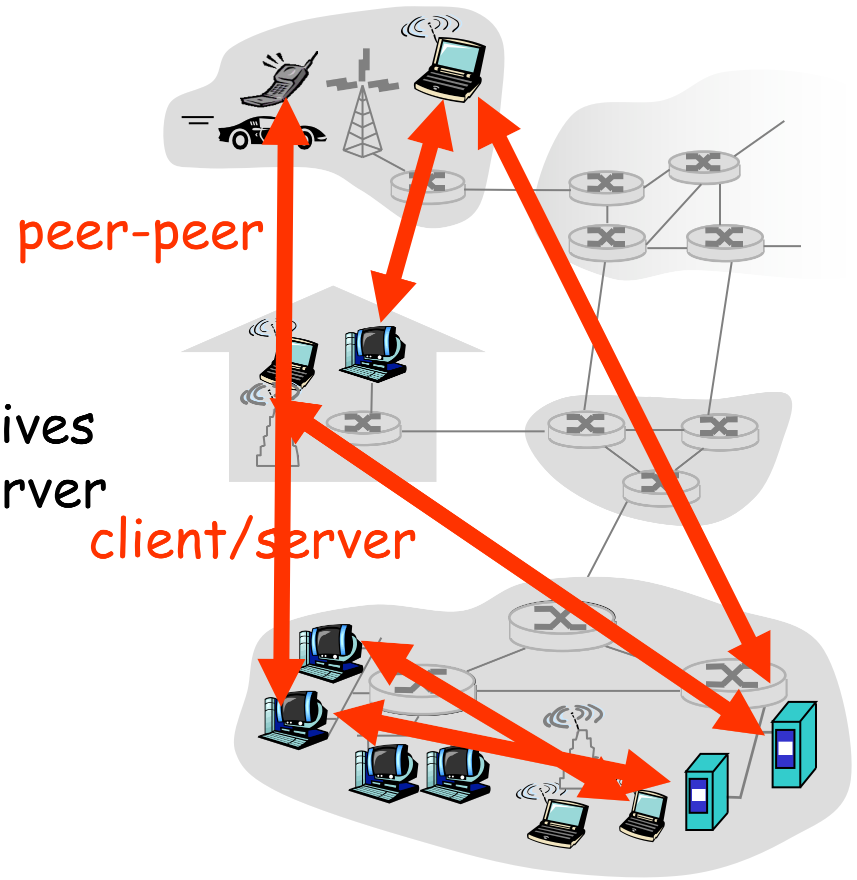

Network edge: connection-oriented service

Goal: data transfer between end systems

- ***Connection:*** prepare for data transfer ahead of time
  - Request / Respond
  - *set up "state"* in two communicating hosts
- TCP - Transmission Control Protocol
  - Internet's connection-oriented service

TCP service [RFC 793]

- *reliable, in-order* byte-stream data transfer
  - loss: acknowledgements and retransmissions
- *flow control:*
  - sender won't overwhelm receiver
- *congestion control:*
  - senders "slow down sending rate" when network congested

Network edge: connectionless service

Goal: data transfer between end systems

- same as before!

- UDP - User Datagram Protocol [RFC 768]:
  - connectionless
  - unreliable data transfer
  - no flow control
  - no congestion control

App's using TCP:

- HTTP (Web), FTP (file transfer), Telnet (remote login), SMTP (email)

App's using UDP:

- streaming media, teleconferencing, DNS, Internet telephony

Chapter 1: roadmap

- What's Computer Network?
- Protocol layers, service models
- What's the Internet?
- Network edge
- Access net and physical media
- Network core
- Internet structure and ISPs
- Delay, loss, and throughput in packet-switched networks
- History of Internet

Access networks and physical media

Q: How to connect end systems to edge router?

- residential access nets
- institutional access networks (school, company)
- mobile access networks

Keep in mind:

- bandwidth of access network?
- shared or dedicated?

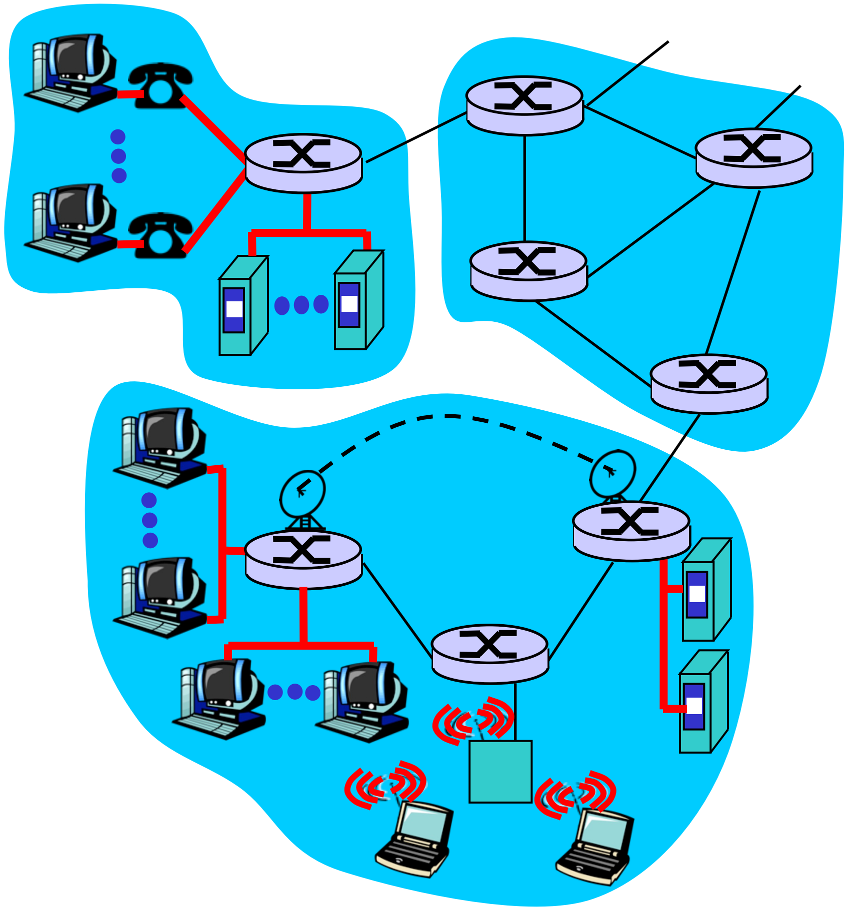

Bandwidth vs Data Rate

- Bandwidth (Hz): the range of frequencies transmitted without being strongly attenuated.

- The bandwidth is a physical property of the transmission medium and usually depends on the construction, thickness, and length of the medium.

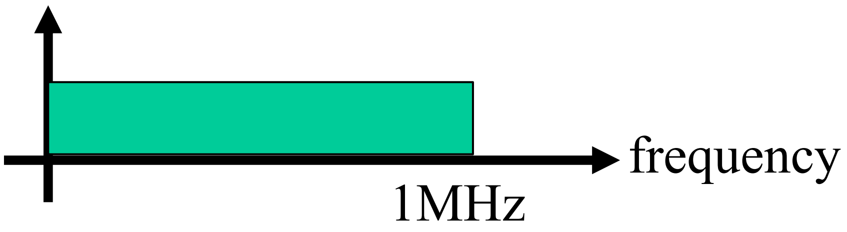

Bandwidth=1MHz

Bandwidth vs Data Rate

- A Wide Band signal will be distorted when transmitted thru relatively narrower band channel with the higher harmonics cut off or hold back.

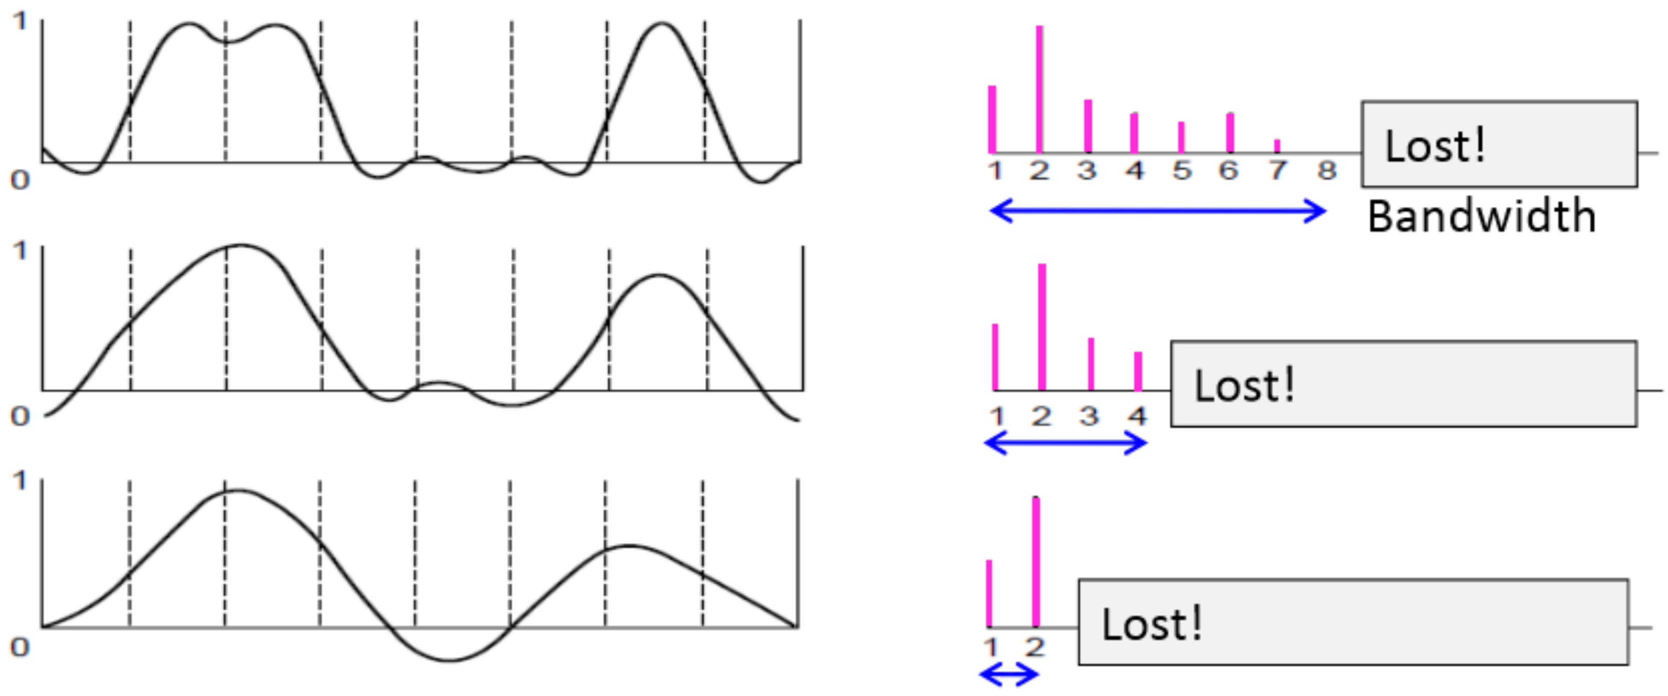

Bandwidth vs Data Rate

- **Data Rate (aka Bit Rate) (bit/s):** the rate at which bits can be transmitted.

- **Shannon's theorem**: the maximum data rate of a noisy channel with signal-to-noise ratio S/N is:

$$R = B\log_2(1+S/N)$$

Data rate (bit/s)

Bandwidth (Hz)

Access networks and physical media

Q: How to connect end systems to edge router?

- residential access nets
- institutional access networks (school, company)
- mobile access networks

Keep in mind:

- bandwidth of access network?
- shared or dedicated?

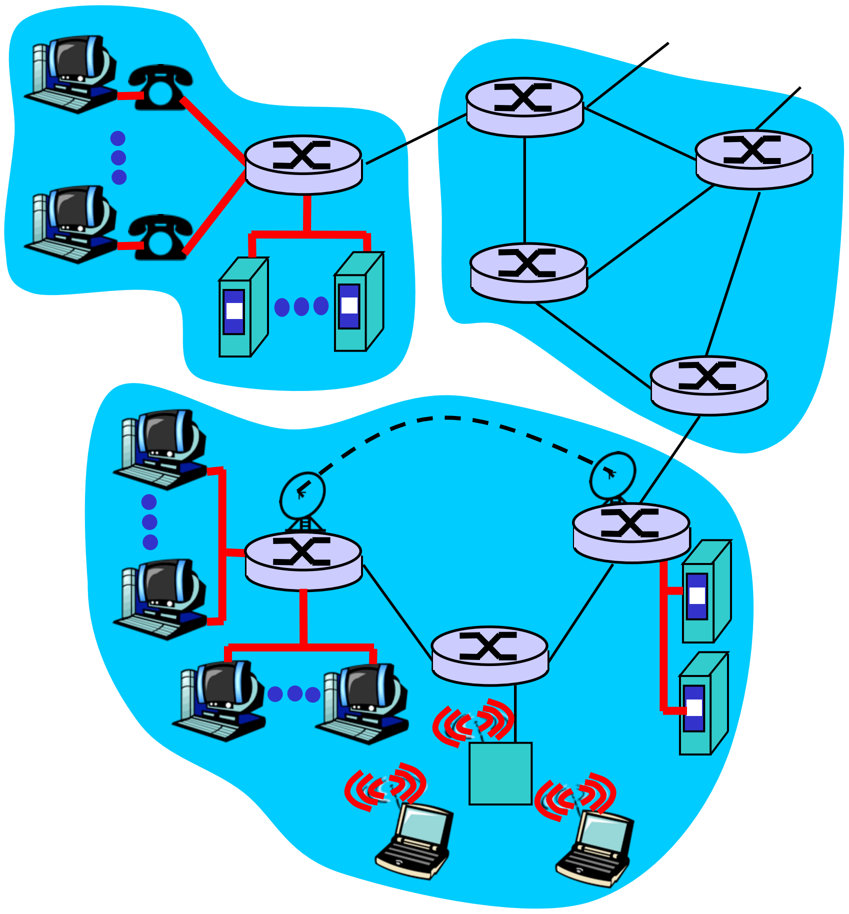

Residential access: point to point access

- Dialup via modem
  - up to 56Kbps direct access to router (often less)
  - Can't surf and phone at same time: can't be "always on"

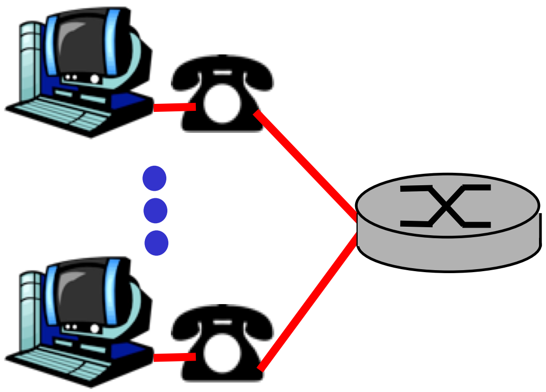

- ADSL: asymmetric digital subscriber line
  - up to 1 Mbps upstream
  - up to 8 Mbps downstream
  - FDM: 50 kHz - 1 MHz for downstream

    4 kHz - 50 kHz for upstream

    0 kHz - 4 kHz for ordinary telephone

Residential access: cable modems

- HFC: hybrid fiber coax
  - asymmetric: up to 30Mbps downstream, 2 Mbps upstream
- network of cable and fiber attaches home to ISP router
  - homes share access to router
- deployment: available via cable TV companies

Residential access: cable modems

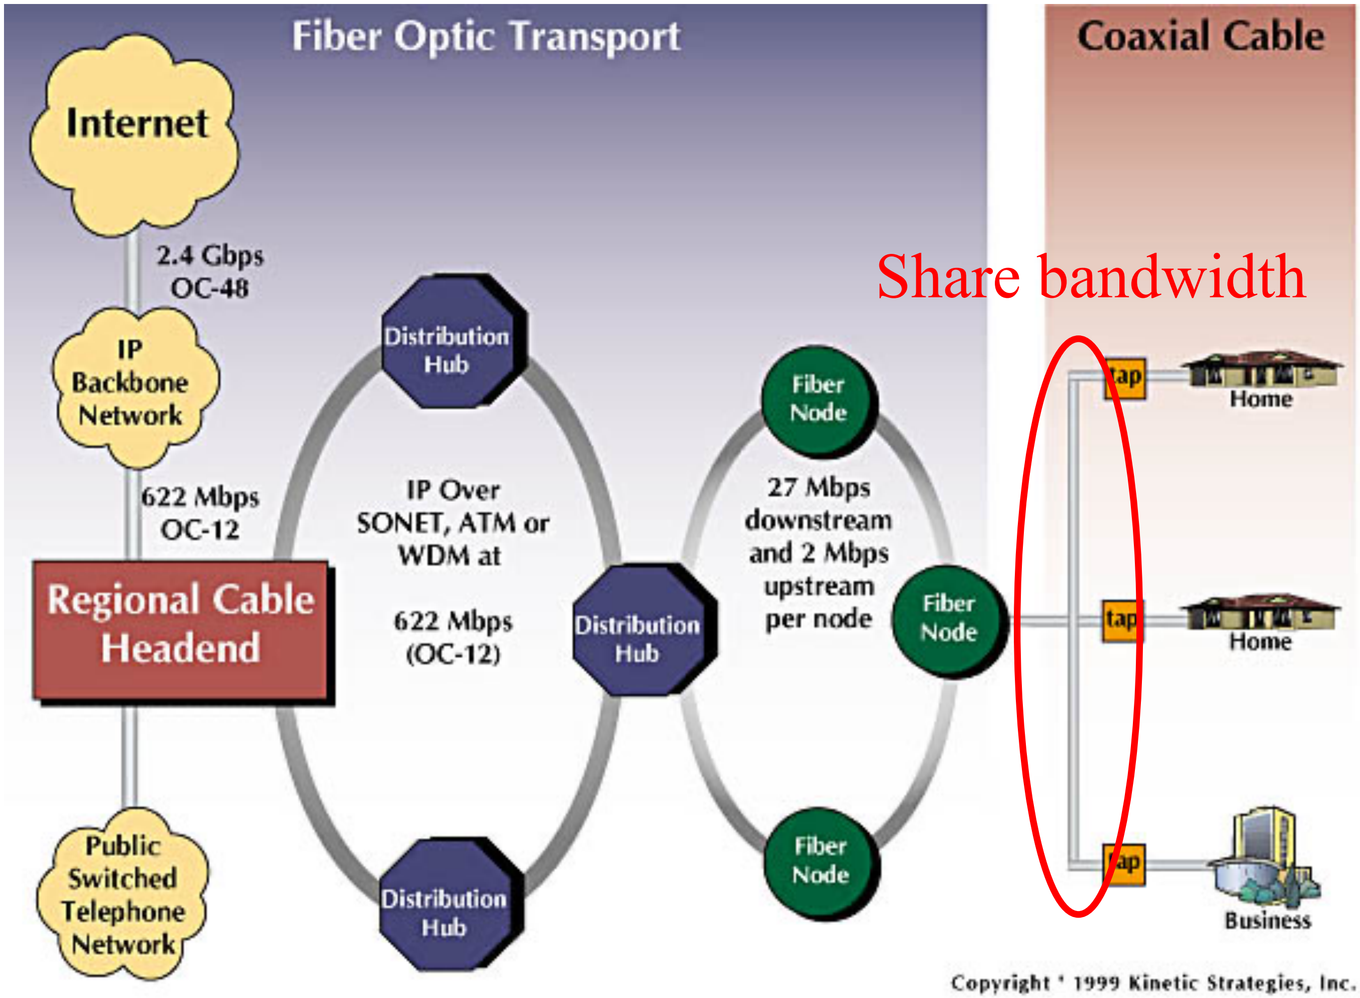

Diagram: http://www.cabledatacomnews.com/cmic/diagram.html

Cable Network Architecture: Overview

Typically 500 to 5,000 homes

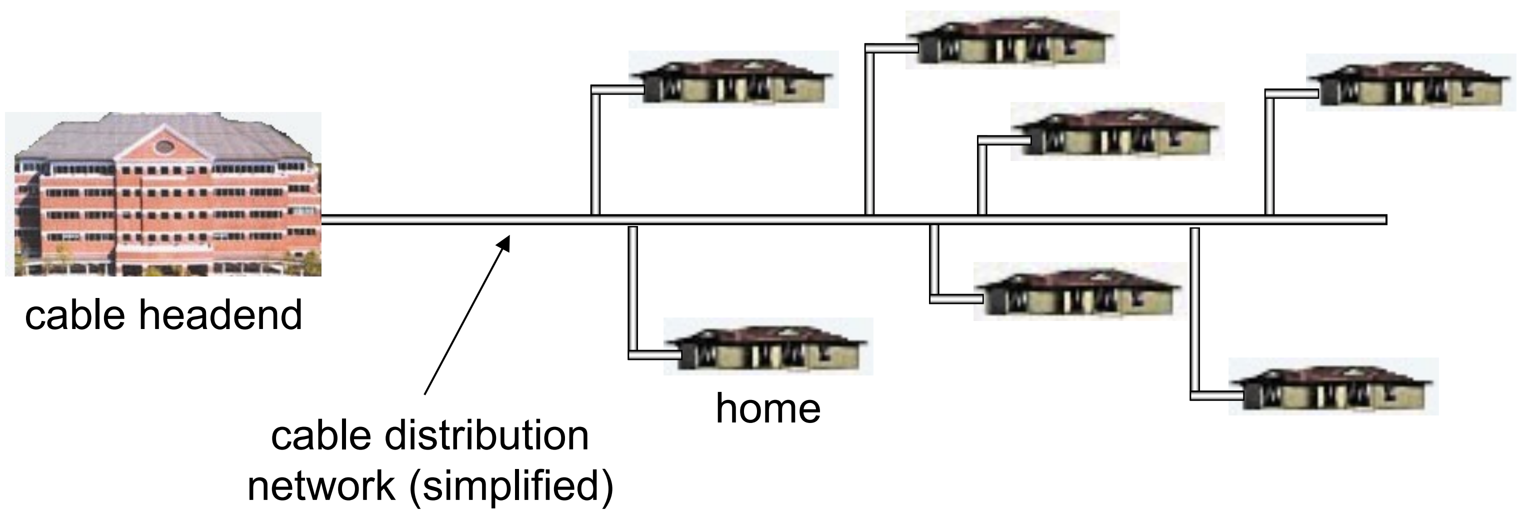

Cable Network Architecture: Overview

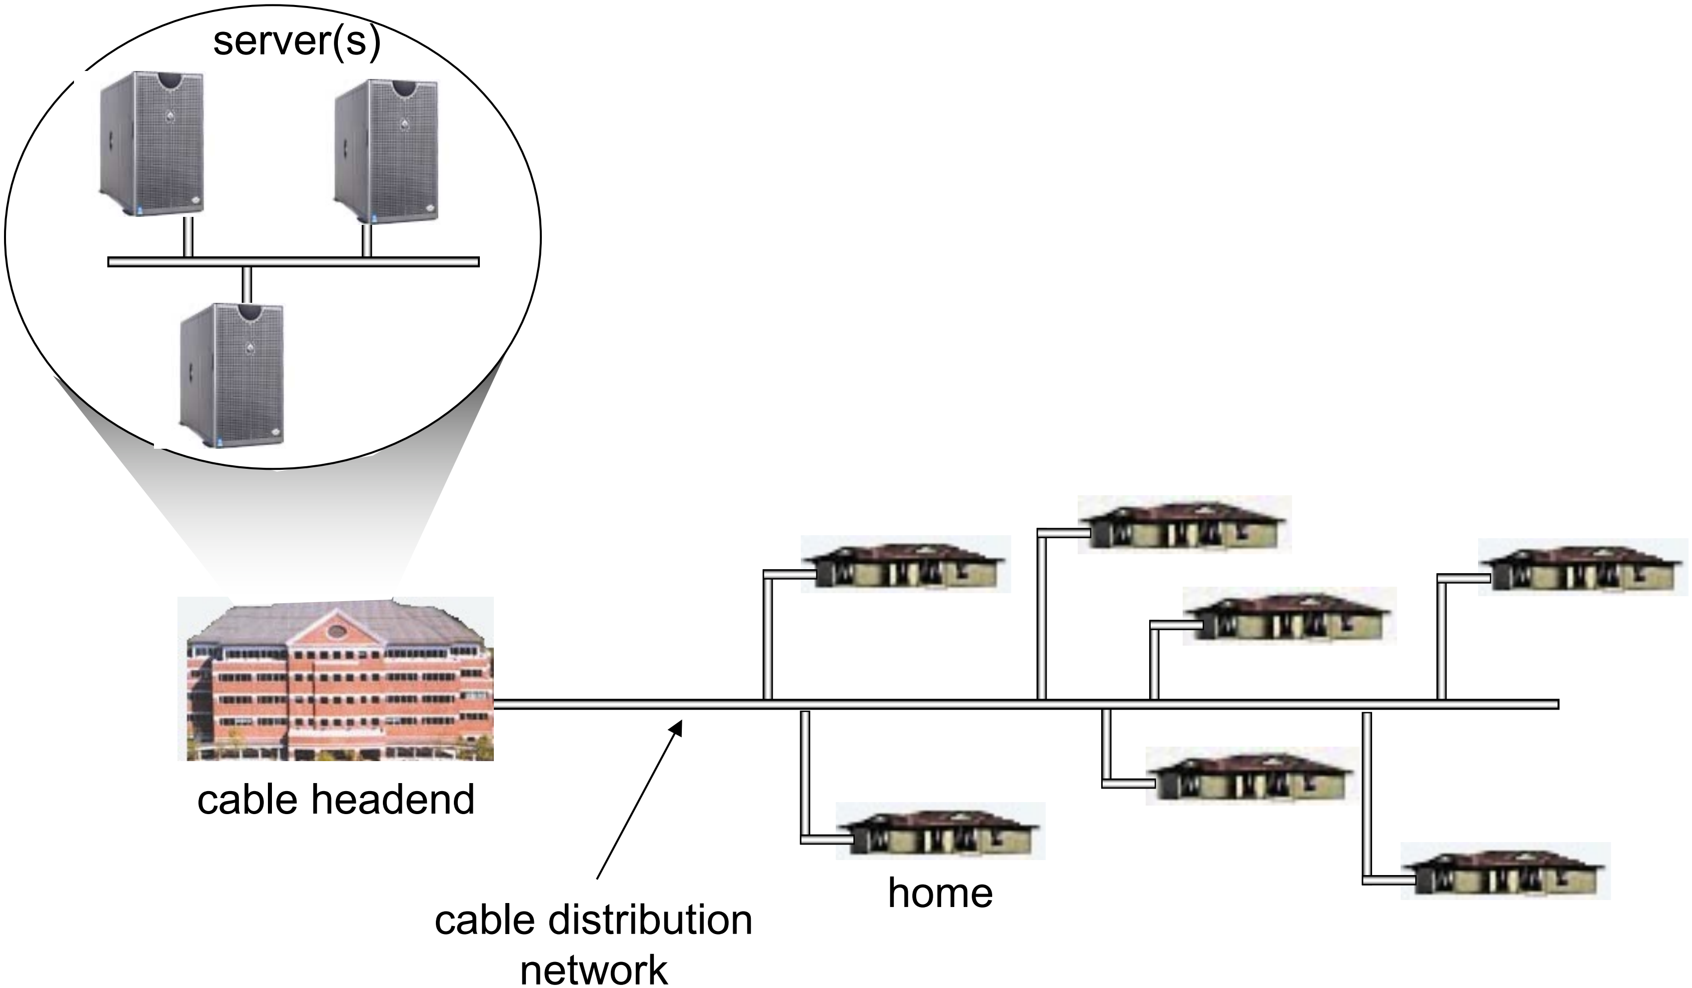

Cable Network Architecture: Overview

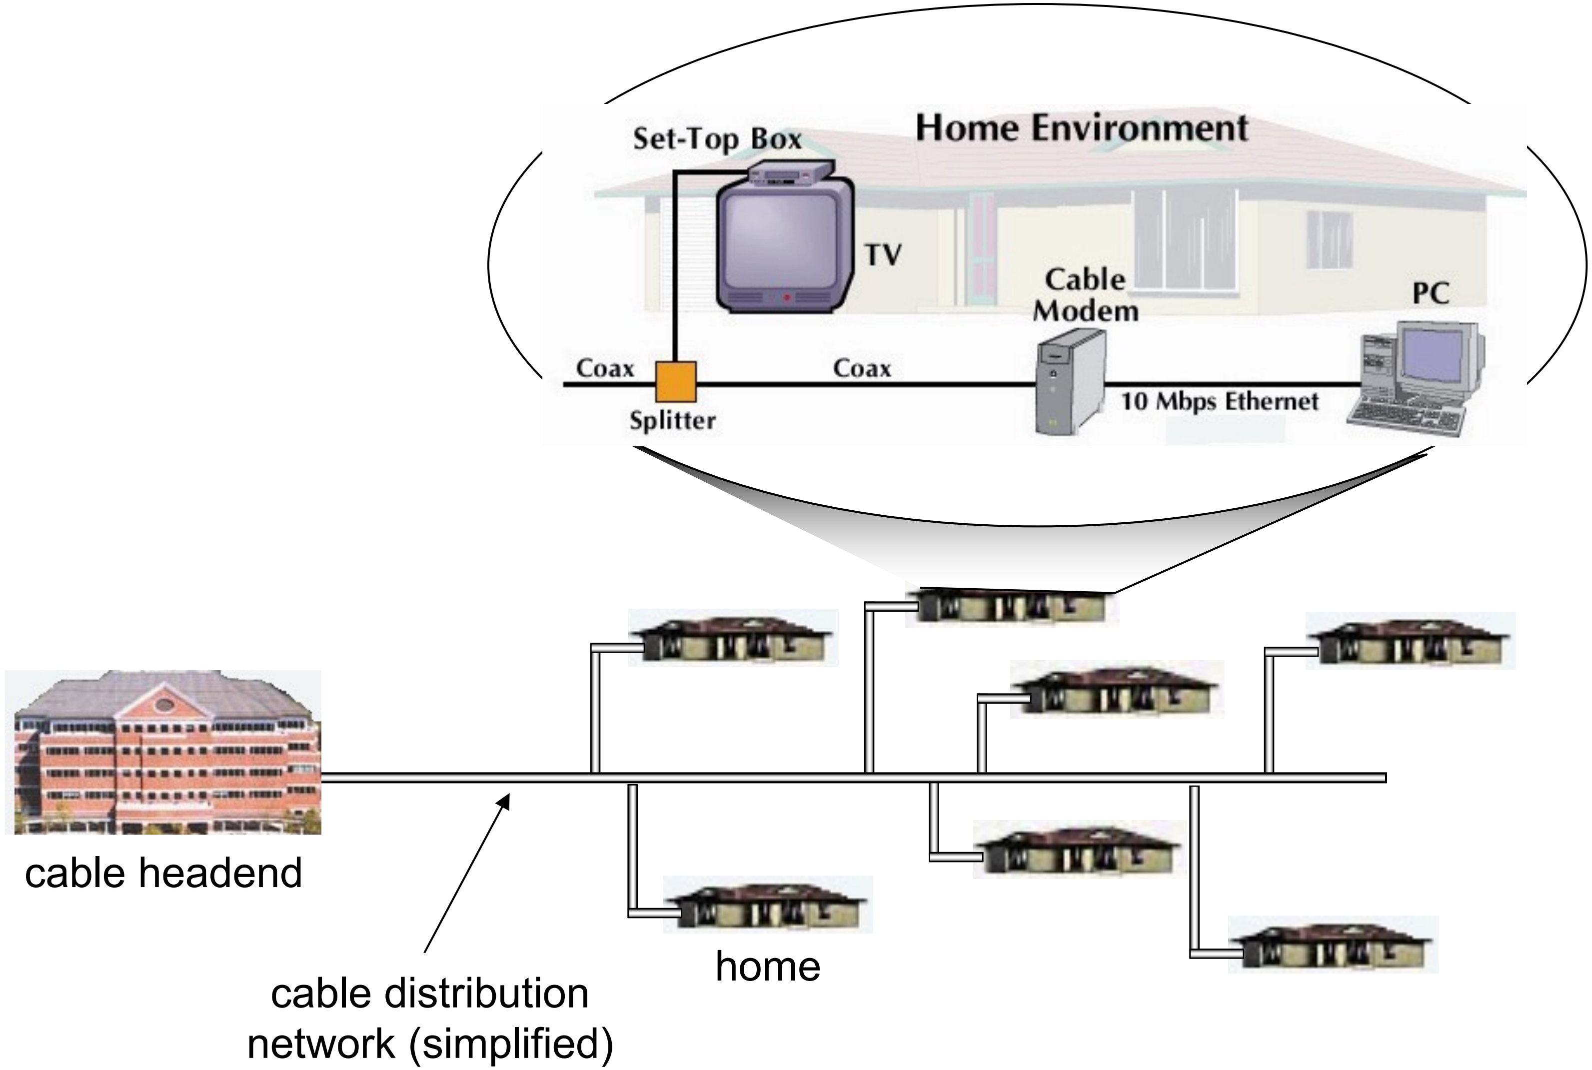

Cable Network Architecture: Overview

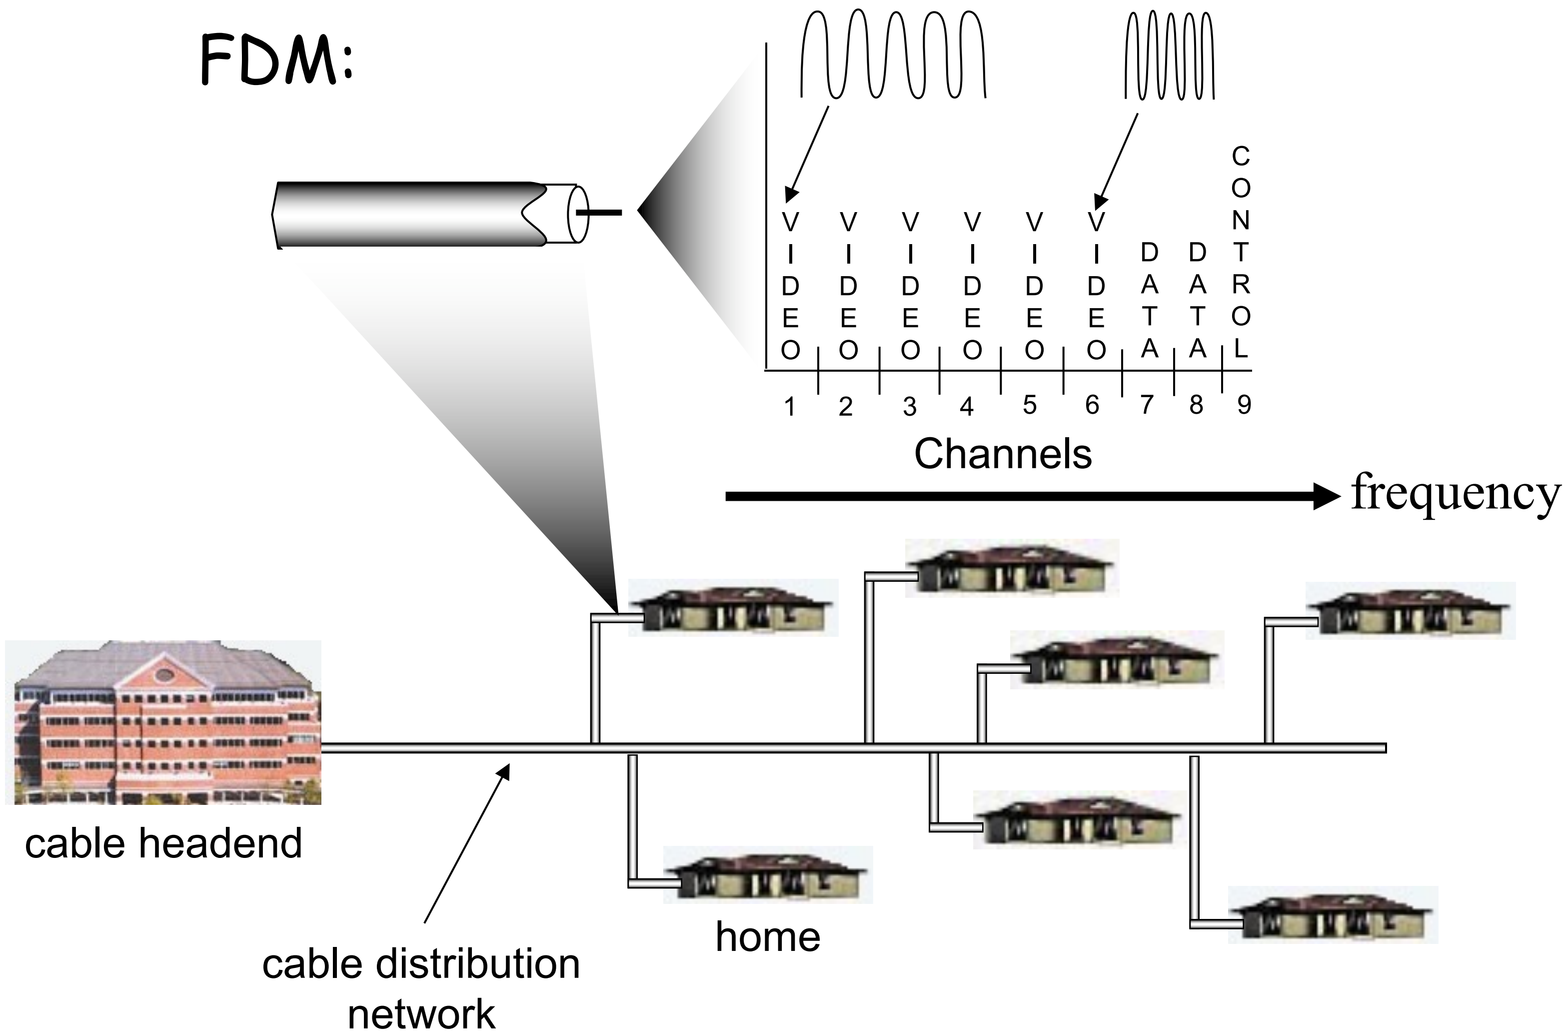

DSL vs Cable Modem

- DSL is point to point

  Thus data rate does not reduce when neighbor uses his/her DSL

- But, DSL uses twisted-pair, and transmission technology cannot support more than ~10Mbps

- Cable Modems share the pipe to the cable headend.

  Thus, your data rate can reduce when neighbors are surfing concurrently

- However, fibre optic lines have significantly higher data rate (fat pipe)

  Even if other users, data rate may still be higher

Company access: local area networks

- company/univ local area network (LAN) connects end system to edge router
- Ethernet:
  - shared or dedicated link connects end system and router
  - 10 Mbs, 100Mbps, Gigabit Ethernet

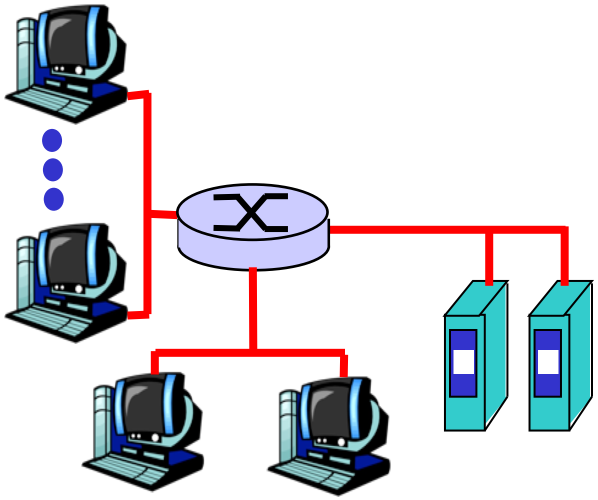

Wireless access networks

- shared *wireless* access network connects end system to router
  - via base station aka "access point"
- wireless LANs:
  - 802.11b/g (WiFi): 11 or 54 Mbps
- wider-area wireless access
  - provided by telco operator

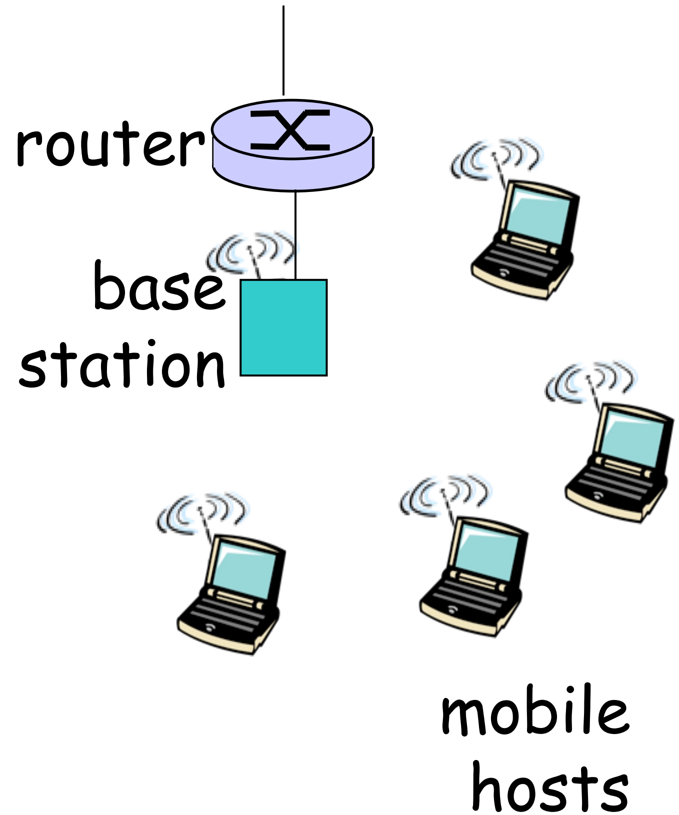

Chapter 1: roadmap

- What's Computer Network?
- Protocol layers, service models
- What's the Internet?
- Network edge
- Access net and physical media
- Network core
- Internet structure and ISPs
- Delay, loss, and throughput in packet-switched networks
- History of Internet

The Network Core

- mesh of interconnected routers
- *the* fundamental question: how is data transferred through net?
  - circuit switching: dedicated circuit per call: telephone net
  - packet-switching: data sent thru net in discrete "chunks"

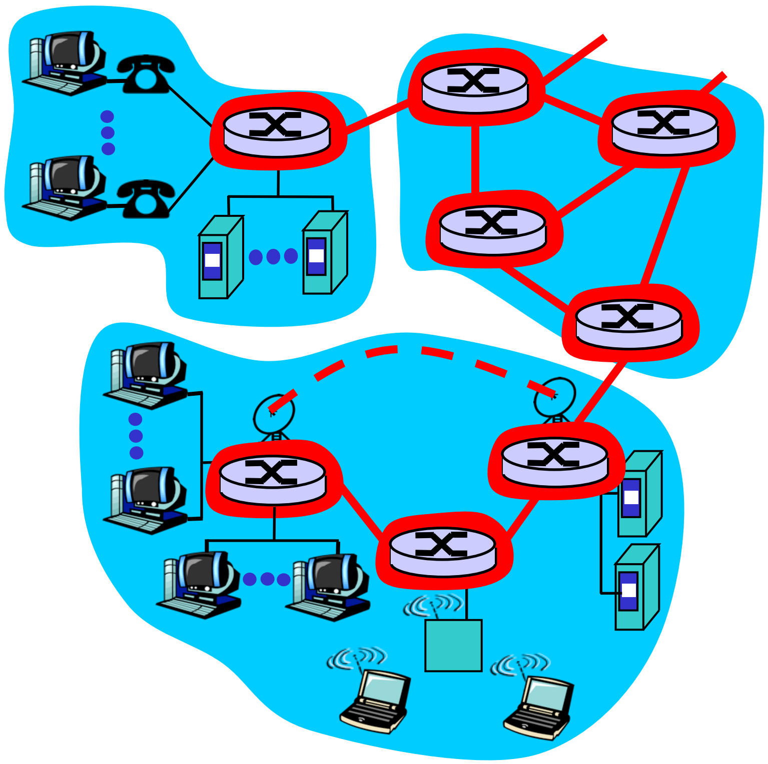

Network Core: Circuit Switching

End-end resources reserved for “call”

- link bandwidth, switch capacity
- dedicated resources: no sharing
- circuit-like (guaranteed) performance
- call setup required

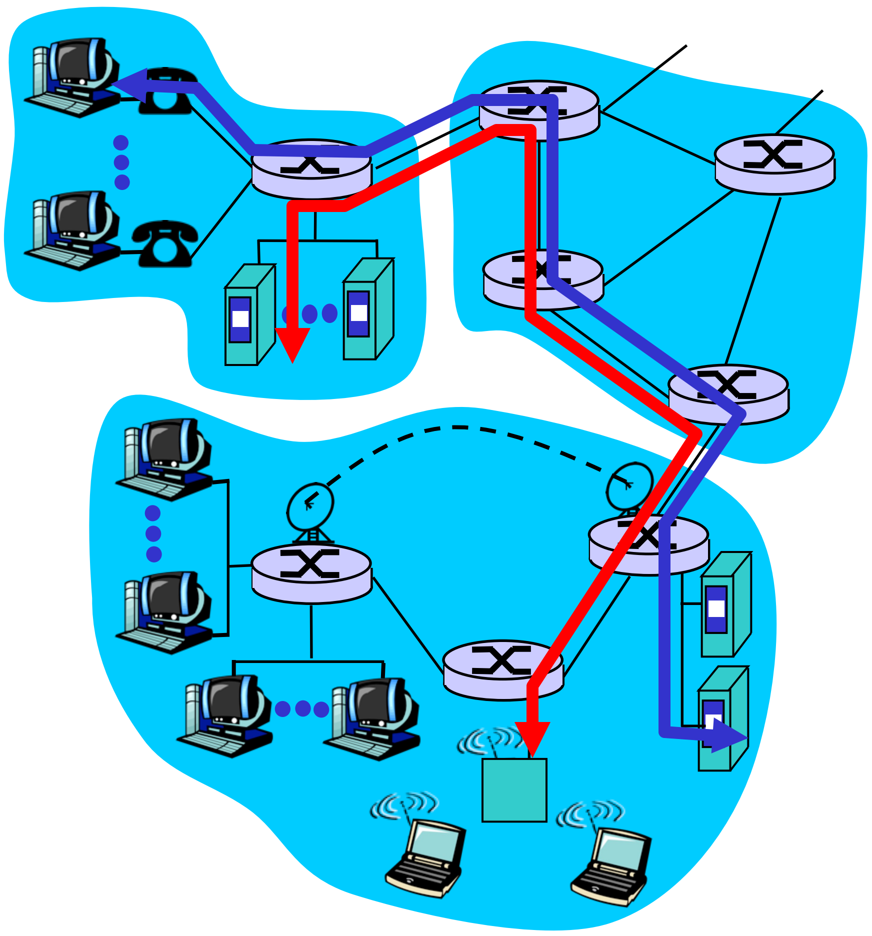

Network Core: Circuit Switching

network resources (e.g., bandwidth) divided into "pieces"

- pieces allocated to calls
- resource piece *idle* if not used by owning call *(no sharing)*
- dividing link bandwidth into "pieces"
  - frequency division
  - time division

Circuit Switching: FDM and TDM

Example:

4 users

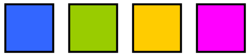

FDM

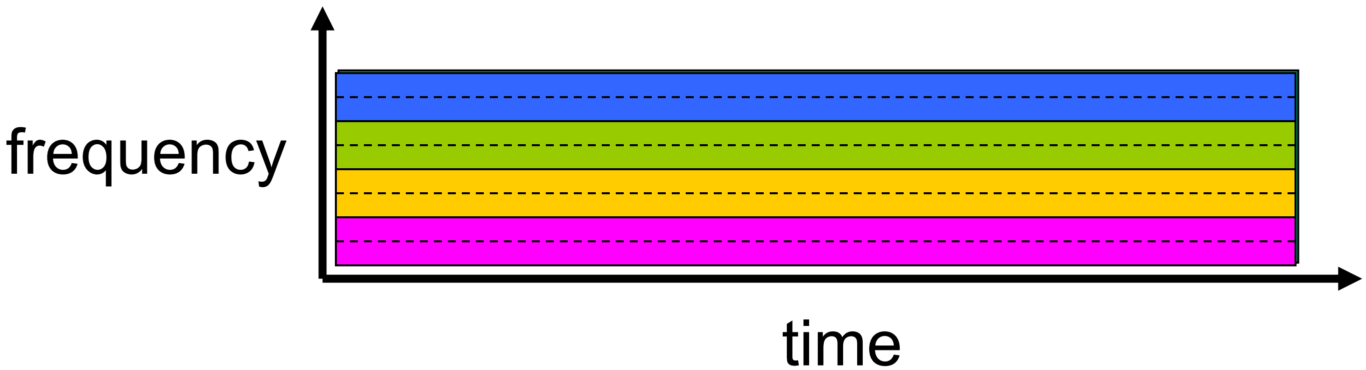

TDM

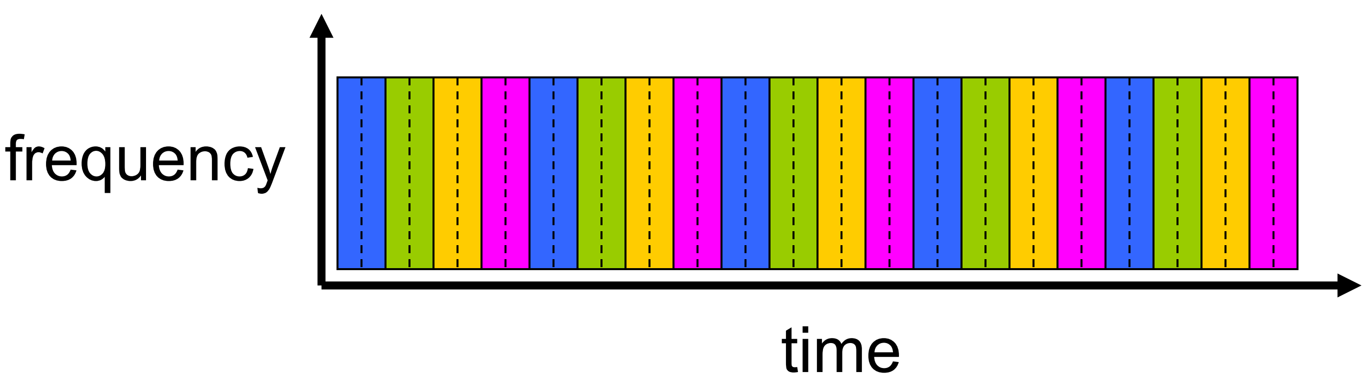

FDM Vs TDM

- What are the tradeoffs?
  - (Dis)Advantage of dividing frequency ?
  - (Dis)Advantage of dividing time ?

Network Core: Packet Switching

each end-end data stream divided into *packets*

- user A, B packets *share* network resources
- each packet uses full link bandwidth
- resources used *as needed*

Bandwidth division into "pieces"

Dedicated allocation

Resource reservation

resource contention:

- aggregate resource demand can exceed amount available
  - Packets queue up
- store and forward: packets move one hop at a time
  - Node receives complete packet before forwarding

Packet Switching: Statistical Multiplexing

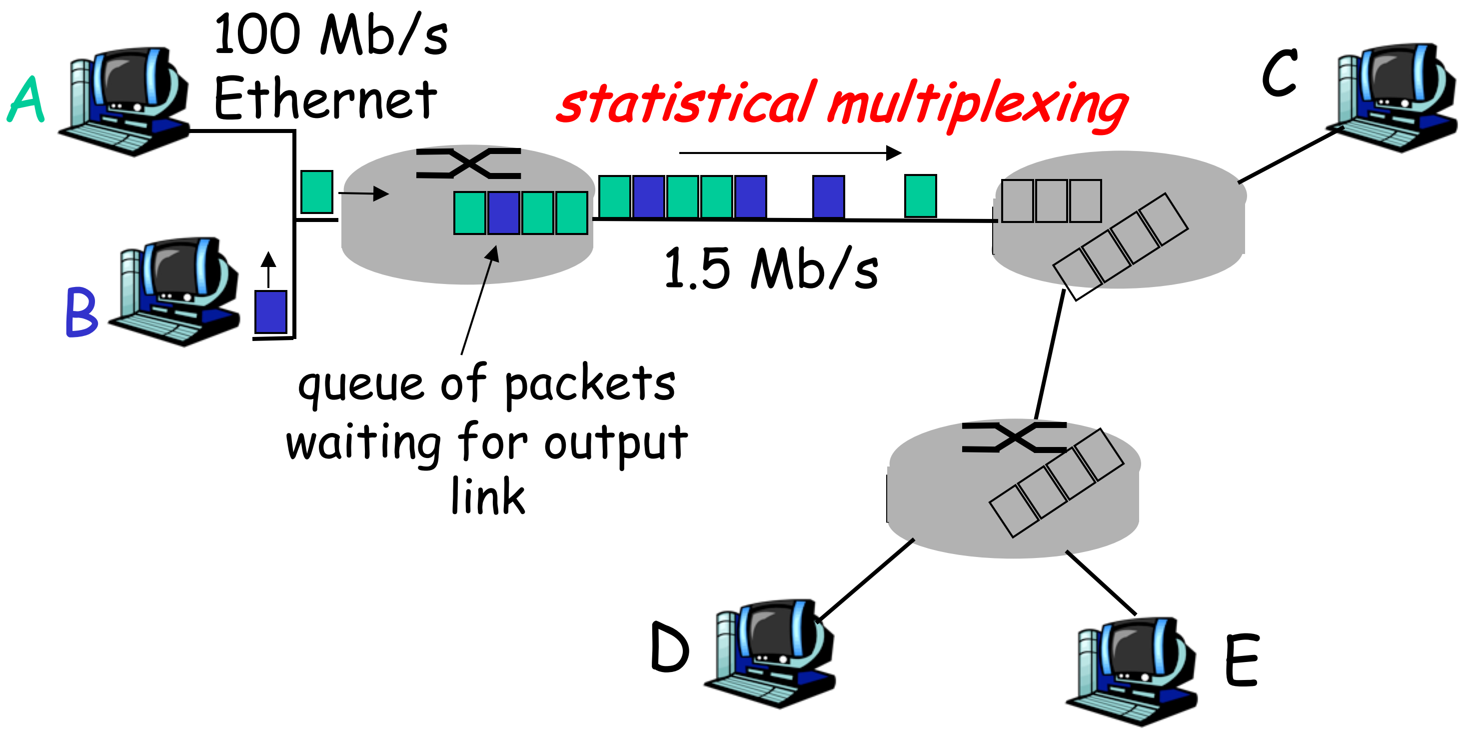

Sequence of A & B packets does not have fixed pattern, shared on demand ➡ ***statistical multiplexing***.

TDM: each host gets same slot in revolving TDM frame.

Compare

Thoughts on tradeoffs between packet switching and circuit switching?

Which one would you take?

Under what circumstances?

Why?

Packet switching versus circuit switching

Packet switching allows more users to use network!

- 1 Mb/s link
- each user:
  - 100 kb/s when "active"
  - active 10% of time

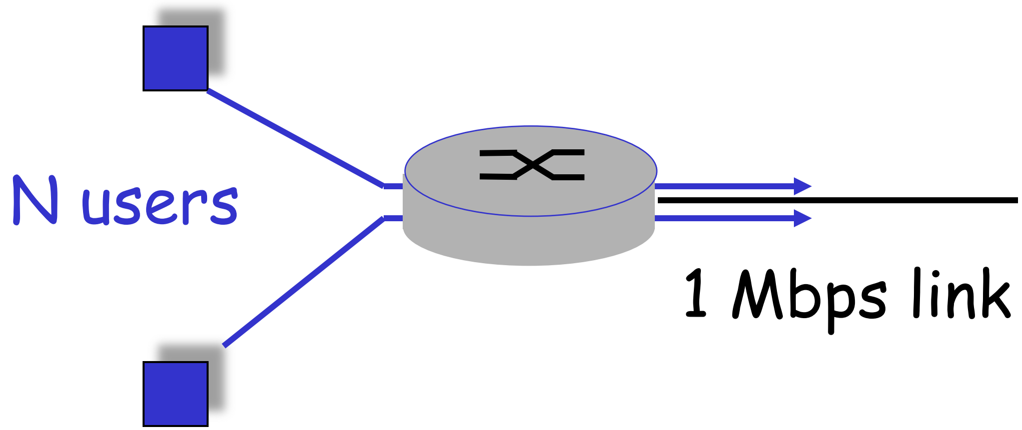

- circuit-switching:
  - 10 users
- packet switching:
  - with 35 users, probability > 10 active less than .0004

Q: how did we get value 0.0004?

Packet switching versus circuit switching

Is packet switching a "slam dunk winner?"

- Great for bursty data
  - resource sharing
  - simpler, no call setup
- Excessive congestion: packet delay and loss
  - protocols needed for reliability, congestion control

Why?

Packet-switching: store-and-forward

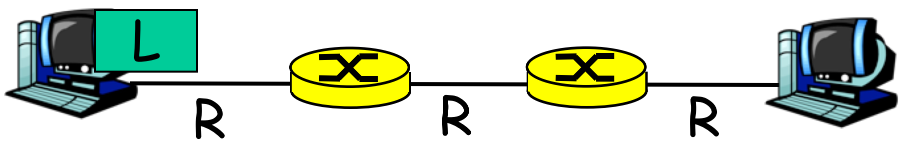

- Takes L/R seconds to transmit (push out) packet of L bits on to link of R bps
- Entire packet must arrive at router before it can be transmitted on next link: *store and forward*
- delay = 3L/R (assuming zero propagation delay)

more on delay shortly …

Example:

- L = 7.5 Mbits
- R = 1.5 Mbps
- delay = 15 sec

Chapter 1: roadmap

- What's Computer Network?
- Protocol layers, service models
- What's the Internet?
- Network edge
- Access net and physical media
- Network core
- Internet structure and ISPs
- Delay, loss, and throughput in packet-switched networks
- History of Internet

Chapter 1: roadmap

- What's Computer Network?
- Protocol layers, service models
- What's the Internet?
- Network edge
- Access net and physical media
- Network core
- Internet structure and ISPs
- Delay, loss, and throughput in packet-switched networks
- History of Internet

How do loss and delay occur?

packets *queue* in router buffers

- packet arrival rate to link (temporarily) exceeds output link capacity
- packets queue, wait for turn

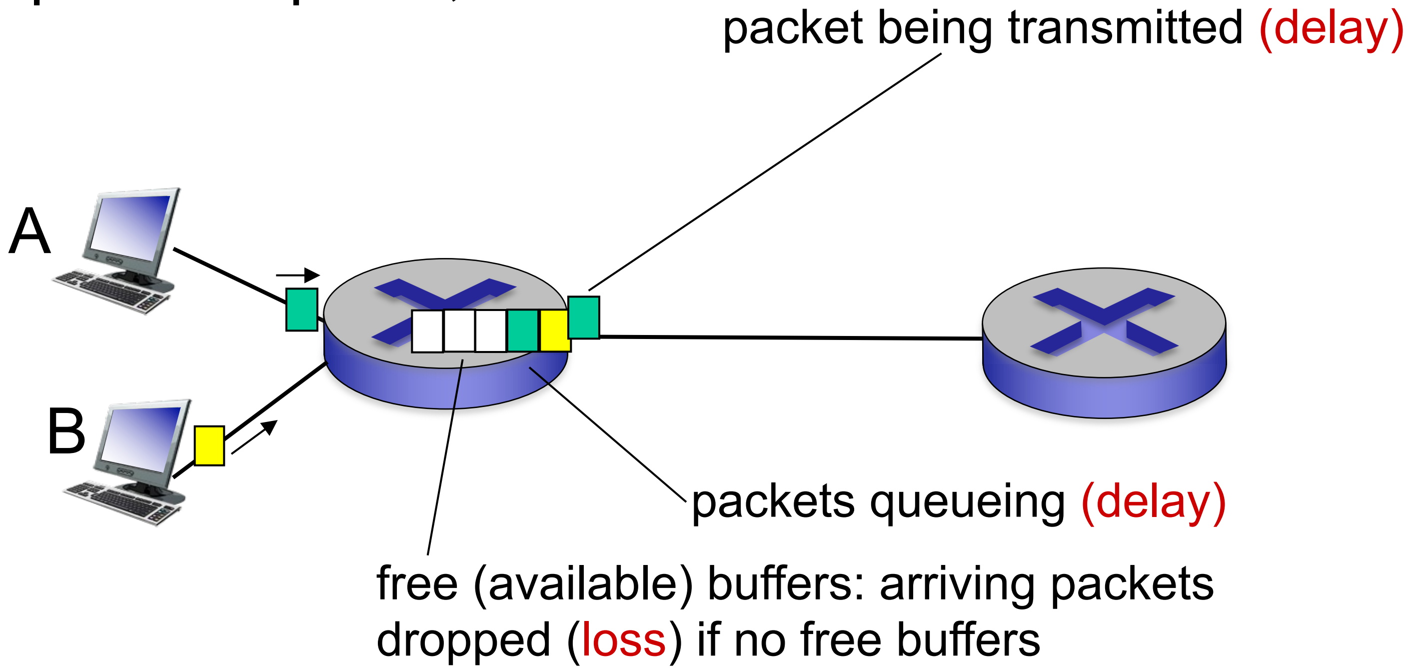

Four sources of packet delay

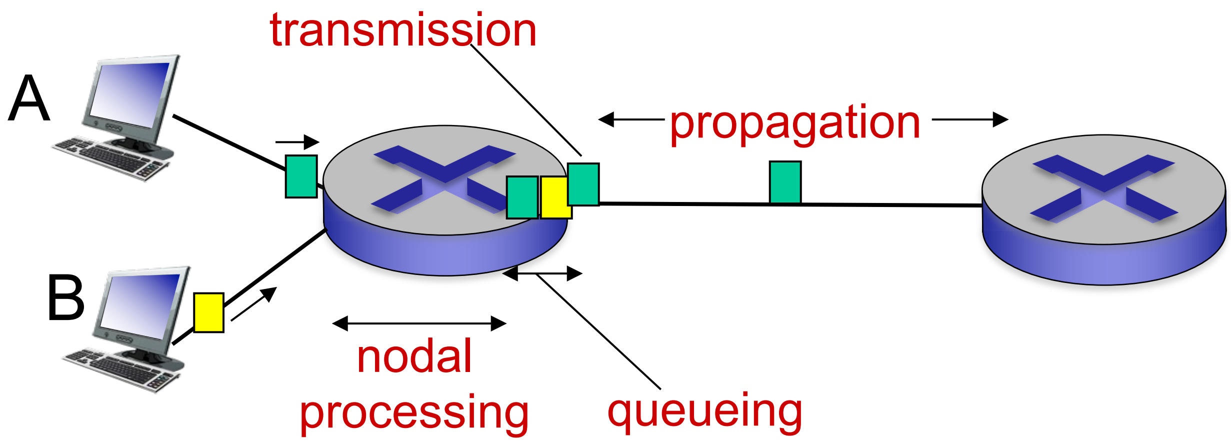

$$d_{nodal} = d_{proc} + d_{queue} + d_{trans} + d_{prop}$$

d_{proc}: nodal processing

- check bit errors
- determine output link
- typically < msec

d_{queue}: queueing delay

- time waiting at output link for transmission
- depends on congestion level of router

Four sources of packet delay

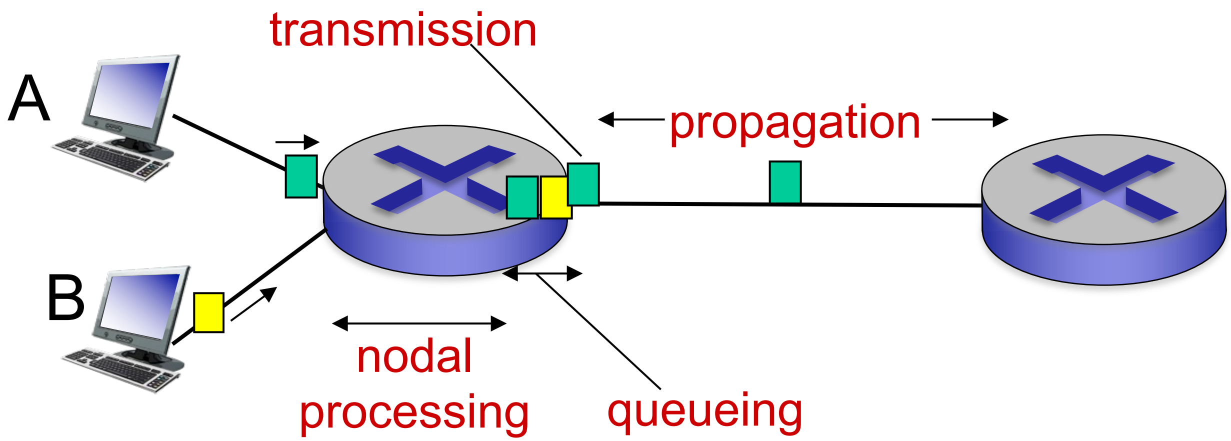

$$d_{nodal} = d_{proc} + d_{queue} + d_{trans} + d_{prop}$$

d_{trans}: transmission delay:
- L: packet length (bits)
- R: link *bandwidth (bps)*
- $d_{trans} = L/R$

d_{trans} and d_{prop} *very* different

d_{prop}: propagation delay:
- d: length of physical link
- s: propagation speed (~2x10^8 m/sec)
- $d_{prop} = d/s$

* Check out the online interactive exercises for more examples: http://gaia.cs.umass.edu/kurose_ross/interactive/

* Check out the Java applet for an interactive animation on trans vs. prop delay

Caravan analogy

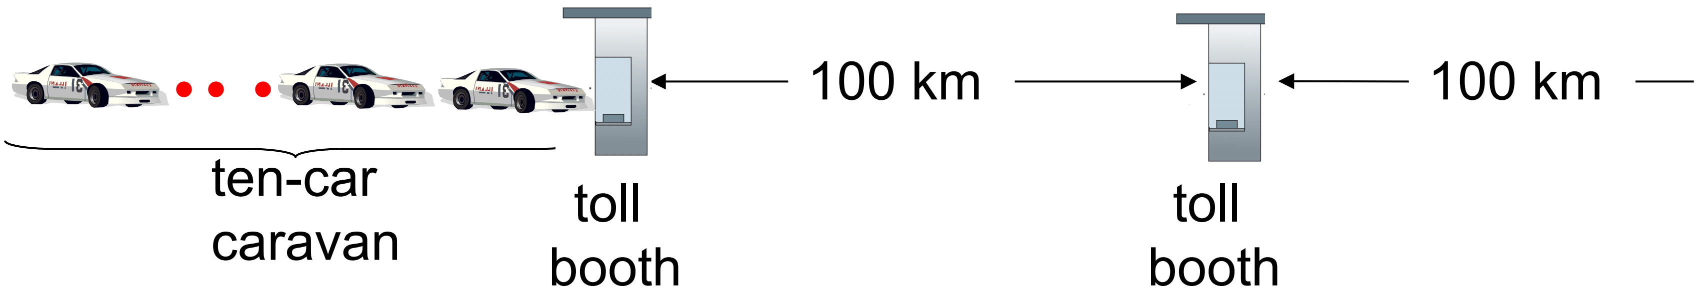

- cars "propagate" at 100 km/hr
- toll booth takes 12 sec to service car (bit transmission time)
- car ~ bit; caravan ~ packet
- *Q:* How long until caravan is lined up before 2nd toll booth?
- time to "push" entire caravan through toll booth onto highway = 12*10 = 120 sec
- time for last car to propagate from 1st to 2nd toll both: 100km/(100km/hr)= 1 hr
- *A:* 62 minutes

Caravan analogy (more)

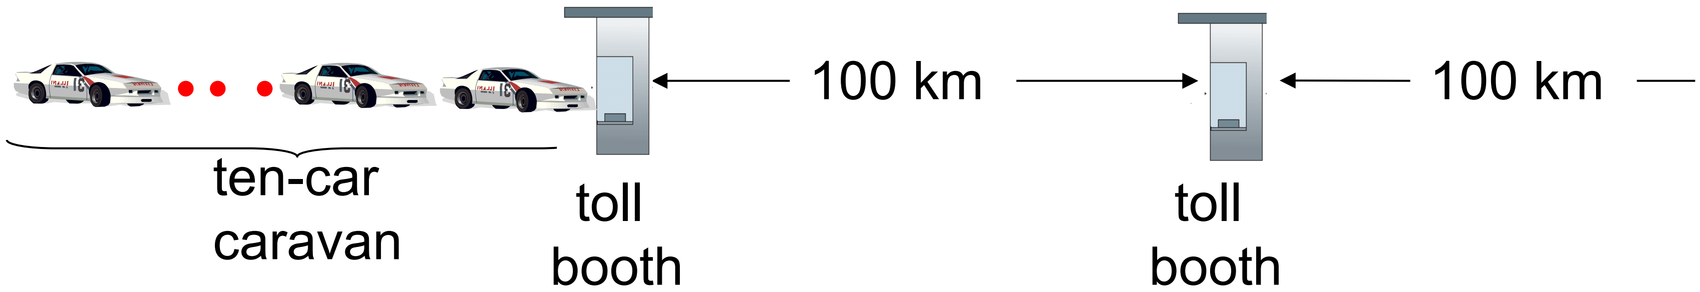

- suppose cars now "propagate" at 1000 km/hr
- and suppose toll booth now takes one min to service a car
- *Q:* Will cars arrive to 2nd booth before all cars serviced at first booth?
  - *A: Yes!* after 7 min, first car arrives at second booth; three cars still at first booth

Queueing delay

- *R:* link bandwidth (bps)
- *L:* packet length (bits)
- a: average packet arrival rate

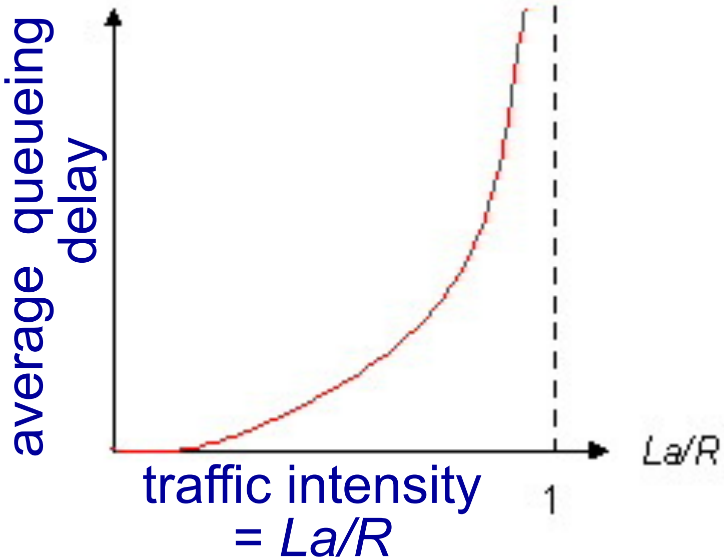

- *La/R* ~ 0: avg. queueing delay small
- *La/R* -> 1: avg. queueing delay large
- "work" arriving than can be serviced, average delay infinite!


La/R ~ 0


La/R -> 1

* Check online interactive animation on queuing and loss

Packet loss

- queue (aka buffer) preceding link in buffer has finite capacity
- packet arriving to full queue dropped (aka lost)
- lost packet may be retransmitted by previous node, by source end system, or not at all

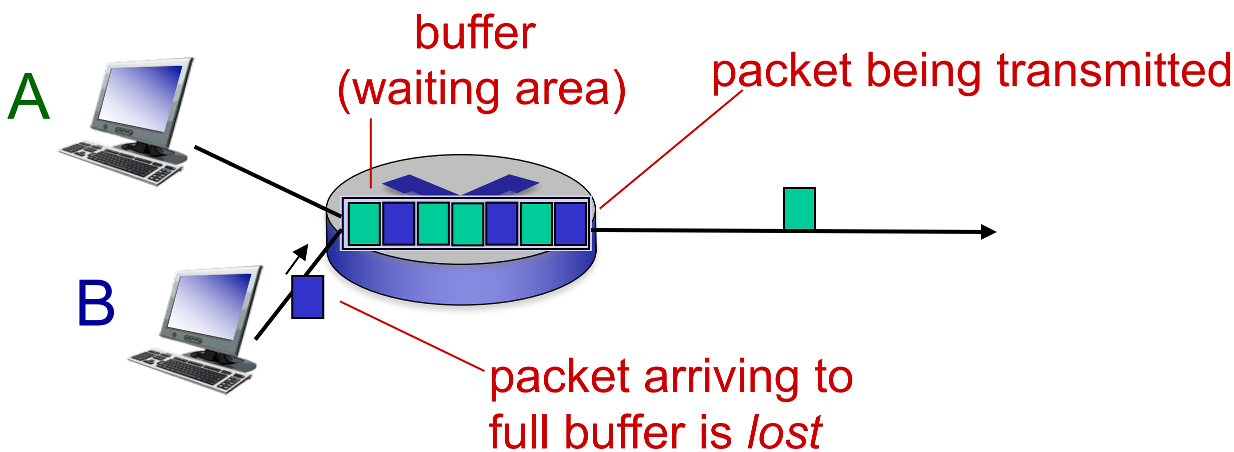

* Check out the Java applet for an interactive animation on queuing and loss

Throughput

- *throughput:* rate (bits/time unit) at which bits transferred between sender/receiver
  - *instantaneous:* rate at given point in time
  - *average:* rate over longer period of time

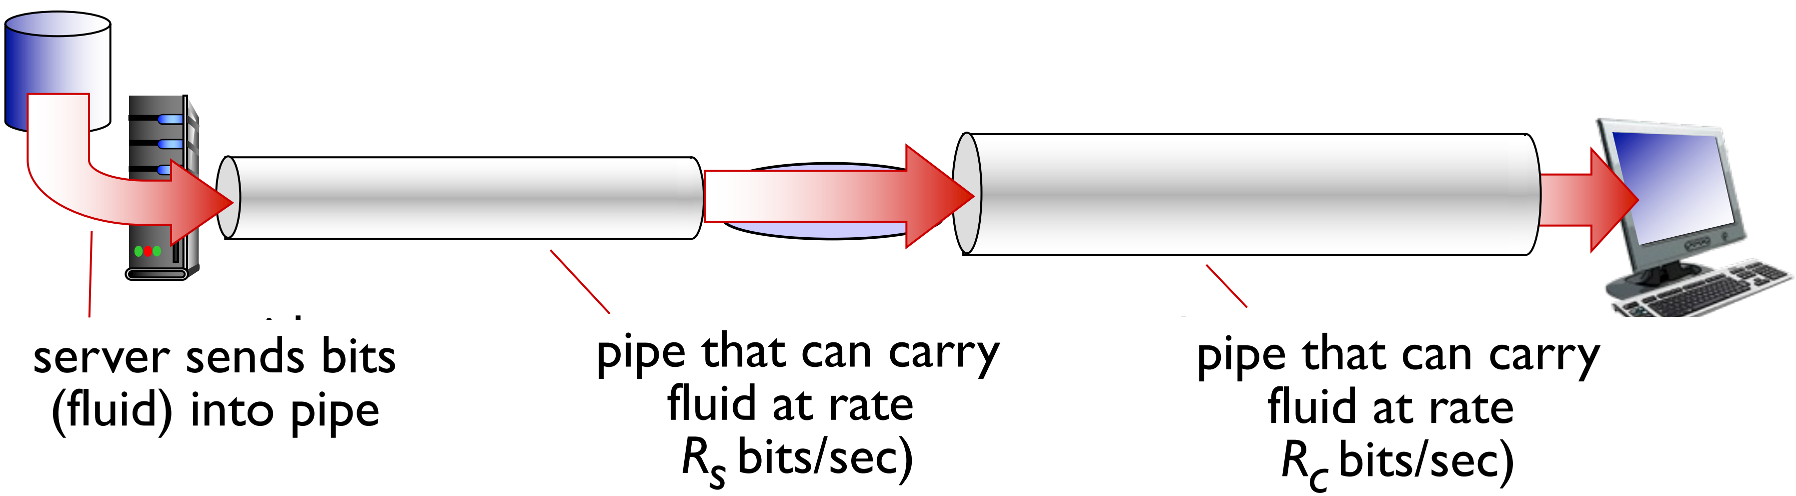

Throughput (more)

- $R_s < R_c$ What is average end-end throughput?

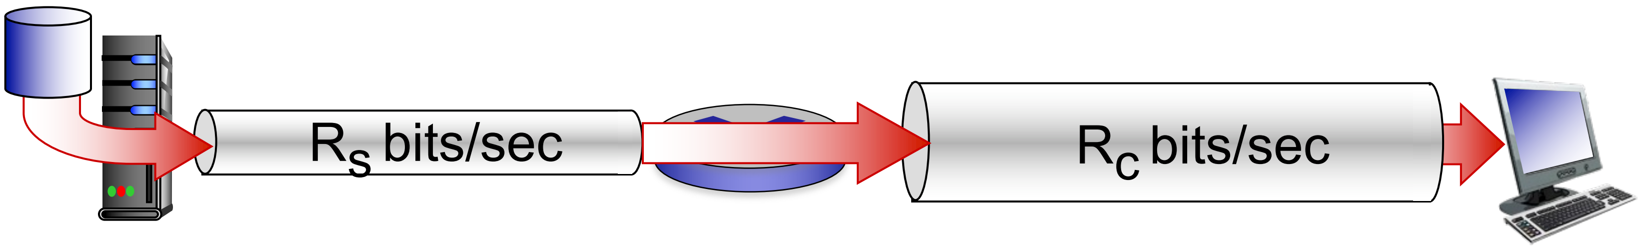

- $R_s > R_c$ What is average end-end throughput?

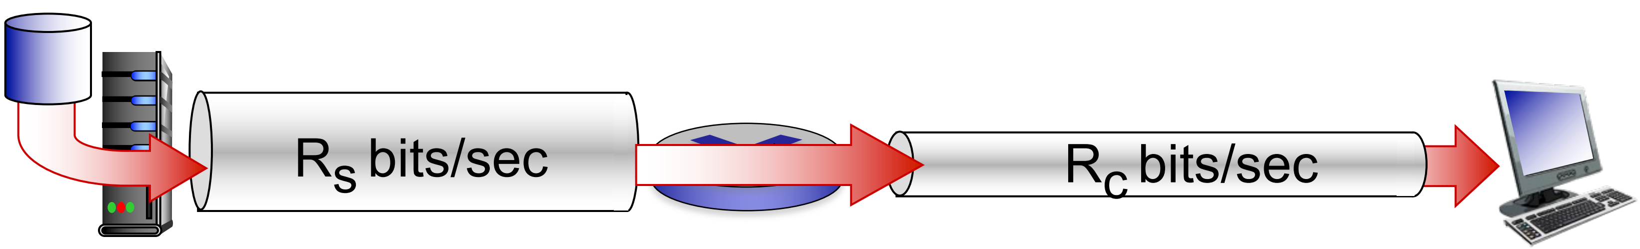

bottleneck link

link on end-end path that constrains end-end throughput

Throughput: Internet scenario

- per-connection end-end throughput: $min(R_c,R_s,R/10)$
- in practice: R_c or R_s is often bottleneck

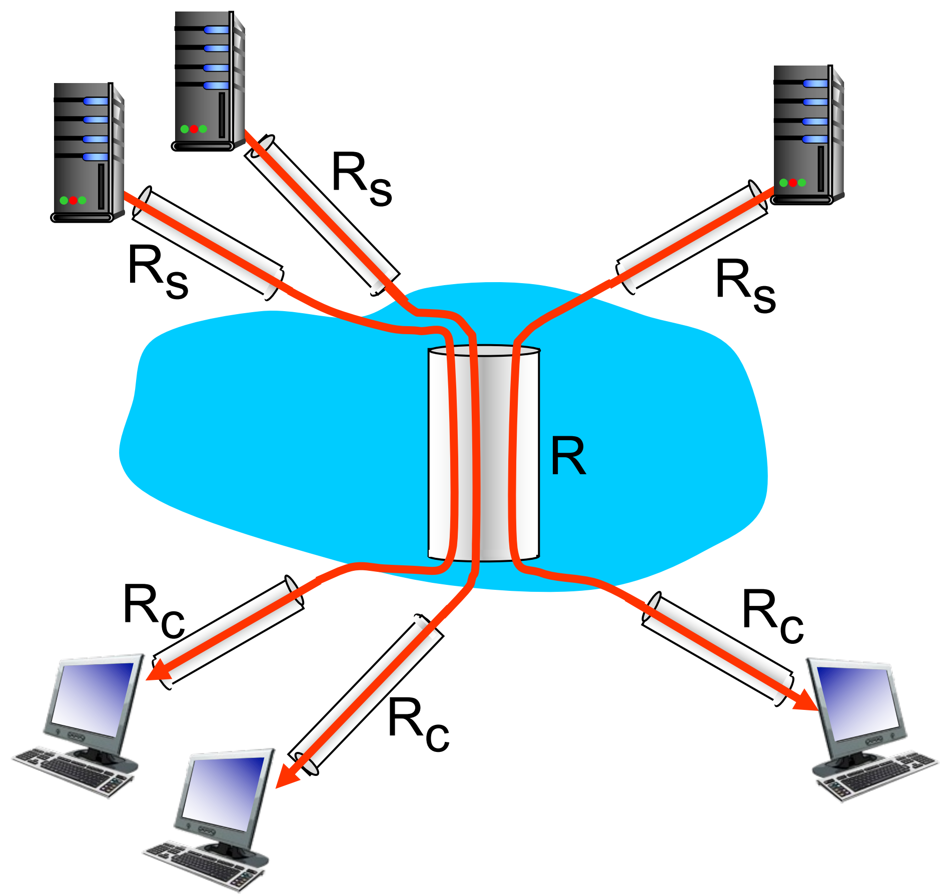

10 connections (fairly) share backbone bottleneck link *R* bits/sec

* Check out the online interactive exercises for more examples: http://gaia.cs.umass.edu/kurose_ross/interactive/

Questions?

Introduction: Summary

Covered a "ton" of material!

- Internet overview
- what's a protocol?
- network edge, core, access network
  - packet-switching versus circuit-switching
- Internet/ISP structure
- performance: loss, delay
- layering and service models

You now have:

- context, overview, "feel" of networking
- more depth, detail *to follow!*

Chapter 1: roadmap

- What's Computer Network?
- Protocol layers, service models
- What's the Internet?
- Network edge
- Access net and physical media
- Network core
- Internet structure and ISPs
- Delay, loss, and throughput in packet-switched networks
- History of Internet

Chapter 1: roadmap

1.1 What *is* the Internet?

1.2 Network edge

1.3 Network core

1.4 Network access and physical media

1.5 Internet structure and ISPs

1.6 Delay & loss in packet-switched networks

1.7 Protocol layers, service models

1.8 History

Internet History

1961-1972: Early packet-switching principles

- 1961: Kleinrock - queueing theory shows effectiveness of packet-switching
- 1964: Baran - packet-switching in military nets
- 1967: ARPAnet conceived by Advanced Research Projects Agency
- 1969: first ARPAnet node operational
- 1972:
  - ARPAnet public demonstration
  - NCP (Network Control Protocol) first host-host protocol
  - first e-mail program
  - ARPAnet has 15 nodes

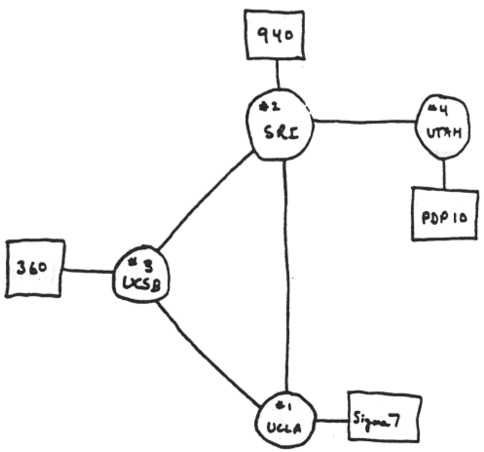

THE ARPA NETWORK

Internet History

1972-1980: Internetworking, new and proprietary nets

- 1970: ALOHAnet satellite network in Hawaii
- 1974: Cerf and Kahn - architecture for interconnecting networks
- 1976: Ethernet at Xerox PARC
- late70's: proprietary architectures: DECnet, SNA, XNA
- late 70's: switching fixed length packets (ATM precursor)
- 1979: ARPAnet has 200 nodes

Cerf and Kahn's internetworking principles:

- minimalism, autonomy - no internal changes required to interconnect networks
- best effort service model
- stateless routers
- decentralized control

define today's Internet architecture

Internet History

1980-1990: new protocols, a proliferation of networks

- 1983: deployment of TCP/IP
- 1982: smtp e-mail protocol defined
- 1983: DNS defined for name-to-IP-address translation
- 1985: ftp protocol defined
- 1988: TCP congestion control
- new national networks: Csnet, BITnet, NSFnet, Minitel
- 100,000 hosts connected to confederation of networks

Internet History

1990, 2000's: commercialization, the Web, new apps

- Early 1990's: ARPAnet decommissioned
- 1991: NSF lifts restrictions on commercial use of NSFnet (decommissioned, 1995)
- early 1990s: Web
  - hypertext [Bush 1945, Nelson 1960's]
  - HTML, HTTP: Berners-Lee
  - 1994: Mosaic, later Netscape
  - late 1990's: commercialization of the Web

Late 1990's – 2000's:

- more killer apps: instant messaging, P2P file sharing
- network security to forefront
- est. 50 million host, 100 million+ users
- backbone links running at Gbps