

1. To add a new default route for vm1,vm2 respectively, we can add following code to step1:

On vm1: sudo route add default gw 192.168.56.201

On vm2: sudo route add default gw 192.168.56.201

2. To Build the VXLAN tunnel and set the remote IP address as 192.168.56.102, add this to vm1 in step2:

```
sudo ovs-vsctl add-port s1 vxlan0 -- set interface vxlan0 type=vxlan
options:remote_ip=192.168.56.102
```

To Build the VXLAN tunnel and set the remote IP address as 192.168.56.101, add this to vm2 in step2:

```
sudo ovs-vsctl add-port s2 vxlan0 -- set interface vxlan0 type=vxlan
options:remote_ip=192.168.56.101
```

2. First check the config

use: sudo ifconfig -a

```
File Edit View Search Terminal Help
[sudo] password for dajiaohuang:
ovs-vsctl: cannot create a bridge named br1 because a bridge named br1 already exists
SIOCADDRT: File exists
dajiaohuang@ubuntu:~/cs3611/lab4$ ifconfig -a
br1: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 192.168.56.101 netmask 255.255.255.0 broadcast 192.168.56.255
    inet6 fe80::ccdf:91ff:fe8:814c prefixlen 64 scopeid 0x20<link>
    ether 00:0c:29:5b:f7:51 txqueuelen 1000 (Ethernet)
    RX packets 500 bytes 39932 (39.9 KB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 530 bytes 44064 (44.0 KB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

ens33: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 192.168.186.133 netmask 255.255.255.0 broadcast 192.168.186.255
    inet6 fe80::3799:7062:cd69:6b68 prefixlen 64 scopeid 0x20<link>
    ether 00:0c:29:5b:f7:51 txqueuelen 1000 (Ethernet)
    RX packets 2740 bytes 612248 (612.2 KB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 1303 bytes 141816 (141.8 KB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
    loop txqueuelen 1000 (Local Loopback)
    RX packets 9789 bytes 770301 (770.3 KB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 9789 bytes 770301 (770.3 KB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

ovs-system: flags=4098<BROADCAST,MULTICAST> mtu 1500
    ether c2:8f:7d:bf:f9:52 txqueuelen 1000 (Ethernet)
    RX packets 0 bytes 0 (0.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

s1: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 10.0.0.3 netmask 255.0.0.0 broadcast 10.255.255.255
    inet6 fe80::3064:5aff:feab:2d4d prefixlen 64 scopeid 0x20<link>
    ether 32:64:5a:ab:2d:4d txqueuelen 1000 (Ethernet)
    RX packets 62 bytes 7644 (7.6 KB)
    RX errors 0 dropped 7 overruns 0 frame 0
    TX packets 53 bytes 7928 (7.9 KB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

s2: flags=4098<BROADCAST,MULTICAST> mtu 1500
    ether ba:93:b7:8b:83:49 txqueuelen 1000 (Ethernet)
    RX packets 0 bytes 0 (0.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

s1-eth1: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet6 fe80::9c7c:e8ff:fe28:294f prefixlen 64 scopeid 0x20<link>
    ether 9e:7c:e8:28:29:4f txqueuelen 1000 (Ethernet)
    RX packets 57 bytes 5066 (5.0 KB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 183 bytes 24361 (24.3 KB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

vxlan svs 4789: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 65500
```

There's no enp0s8, so I modify "enp0s8" in the given code into "ens33", so that there won't be an error. And put those commands in lan1.sh and lan2.sh

```
> 计网 > lab4 > $ lan1.sh
1 |sudo ifconfig s1 10.0.0.3/8 up
2 |sudo ovs-vsctl add-br br1
3 |sudo ovs-vsctl add-port br1 ens33
4 |sudo ifconfig br1 192.168.56.101/24 up
5 |sudo route add default gw 192.168.56.201
6 |sudo ovs-vsctl add-port s1 vxlan0 -- set interface vxlan0 type=vxlan options:remote_ip=
```

Run the sh files.


```
File Edit View Search Terminal Help
dajiaohuang@ubuntu:~/cs3611/lab4$ sudo ./lan1.sh
[sudo] password for dajiaohuang:
ovs-vsctl: cannot create a bridge named br1 because a bridge named br1 already exists
SIOCADDRT: File exists
```

Then use xterm and ping 10.0.0.1 from 10.0.0.2.

```
"Node: h2"
root@ubuntu:~/cs3611/lab4# ping 10.0.0.1
PING 10.0.0.1 (10.0.0.1) 56(84) bytes of data:
64 bytes from 10.0.0.1: icmp_seq=1 ttl=64 time=29.0 ms
64 bytes from 10.0.0.1: icmp_seq=2 ttl=64 time=26.0 ms
64 bytes from 10.0.0.1: icmp_seq=3 ttl=64 time=21.7 ms
64 bytes from 10.0.0.1: icmp_seq=4 ttl=64 time=22.0 ms
64 bytes from 10.0.0.1: icmp_seq=5 ttl=64 time=21.8 ms
64 bytes from 10.0.0.1: icmp_seq=6 ttl=64 time=21.4 ms
64 bytes from 10.0.0.1: icmp_seq=7 ttl=64 time=21.5 ms
64 bytes from 10.0.0.1: icmp_seq=8 ttl=64 time=21.4 ms
64 bytes from 10.0.0.1: icmp_seq=9 ttl=64 time=21.6 ms
^C

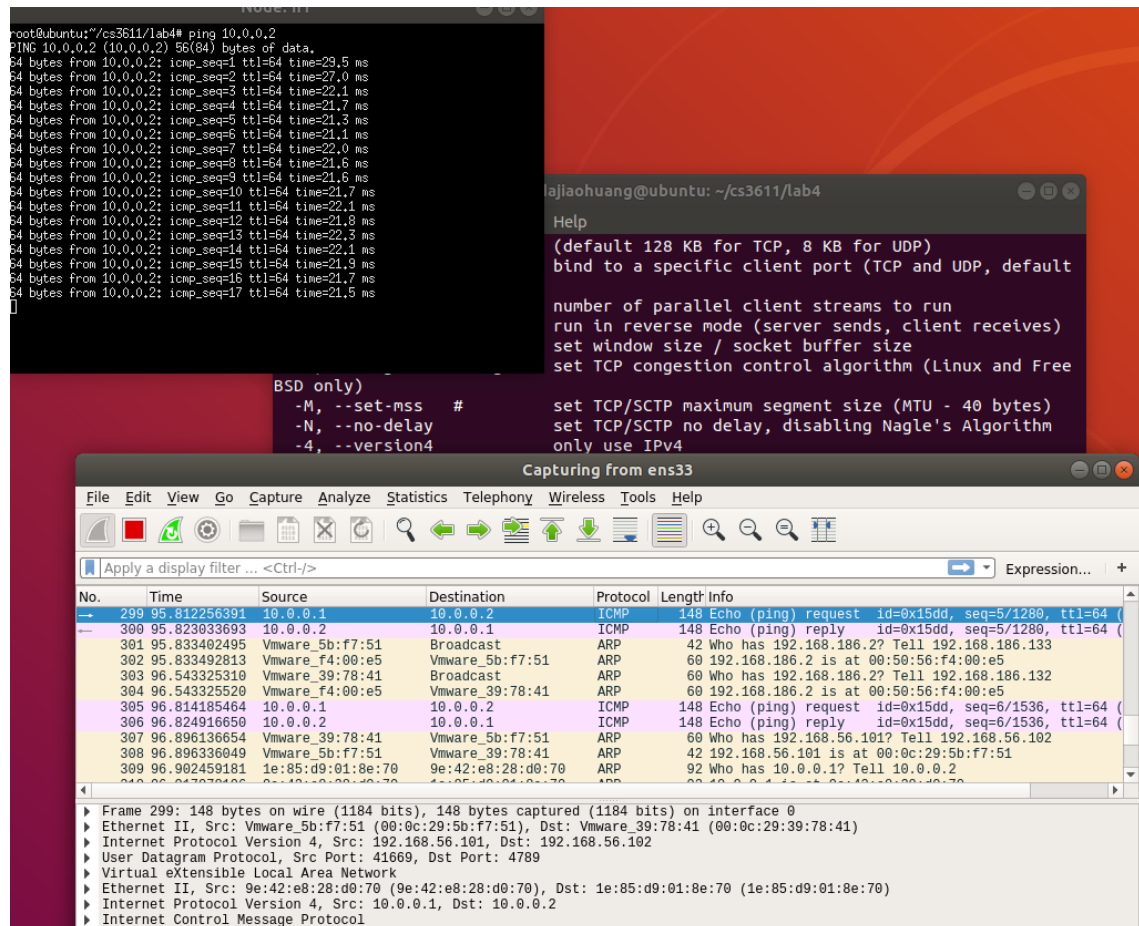
```

Also ping reversely.

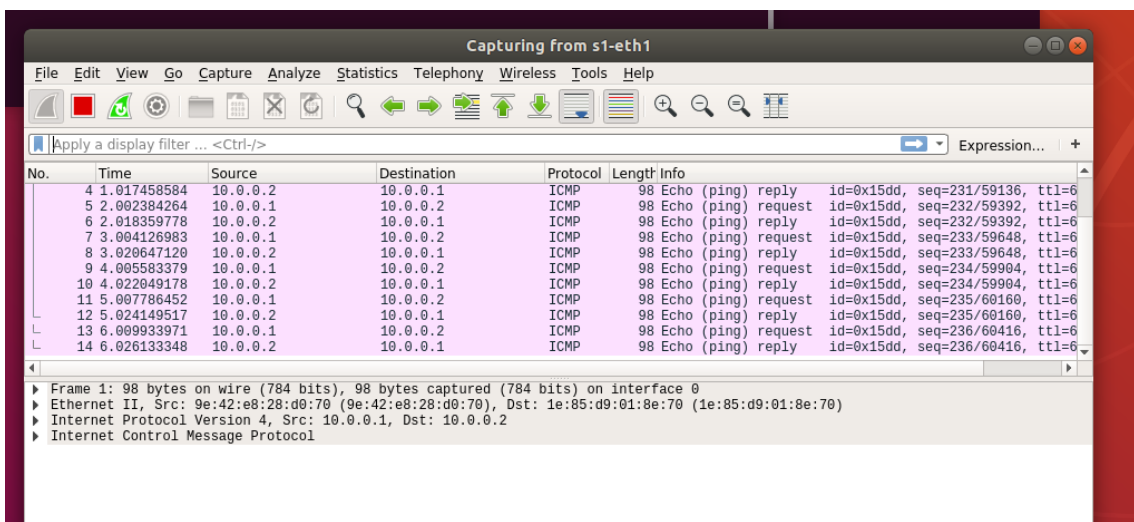


Wireshark when ping 10.0.0.2 from 10.0.0.1:

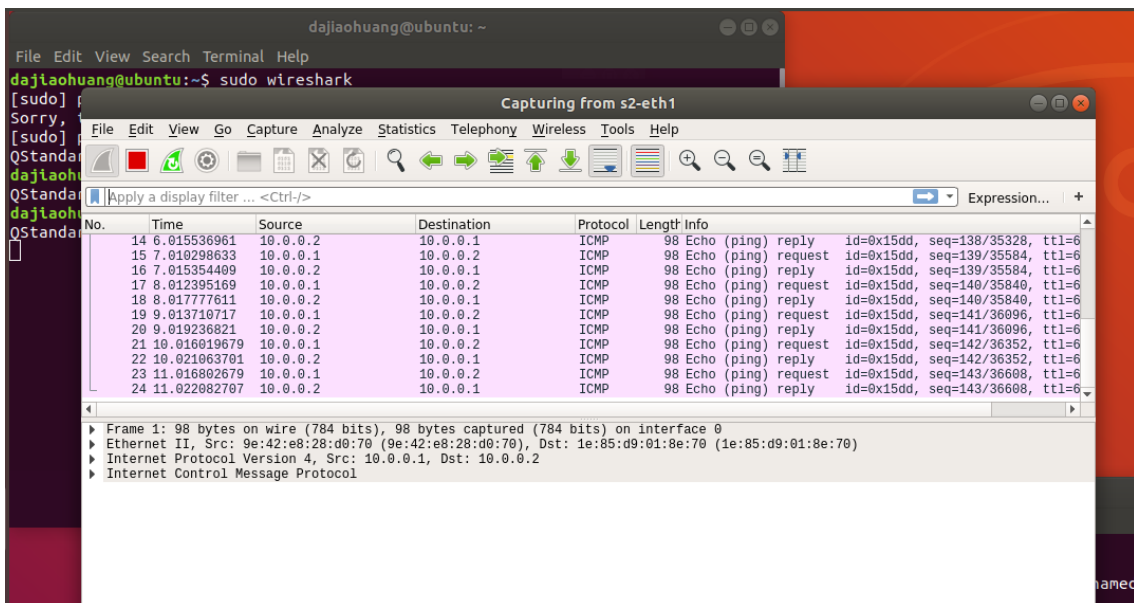
ens33:



s1-eth1:




s2-eth1:



We can see that ICMP,AMP are used.

3. sudo iperf3 -s 10.0.0.4 on s1



- sudo iperf3 -c 10.0.0.3 on s2

```

"Node: s2" (root)
root@ubuntu:~/cs3611/lab4# iperf3 -c 10.0.0.3
Connecting to host 10.0.0.3, port 5201
[ 12] local 10.0.0.4 port 44998 connected to 10.0.0.3 port 5201
[ ID] Interval           Transfer     Bandwidth       Retr   Cwnd
[ 12] 0.00-1.00   sec    76.4 KBytes  625 Kbits/sec    2    1.41 KBytes
[ 12] 1.00-2.00   sec     0.00 Bytes    0.00 bits/sec    1    1.41 KBytes
[ 12] 2.00-3.00   sec     0.00 Bytes    0.00 bits/sec    0    1.41 KBytes
[ 12] 3.00-4.00   sec     0.00 Bytes    0.00 bits/sec    1    1.41 KBytes
[ 12] 4.00-5.00   sec     0.00 Bytes    0.00 bits/sec    0    1.41 KBytes
[ 12] 5.00-6.00   sec     0.00 Bytes    0.00 bits/sec    0    1.41 KBytes
[ 12] 6.00-7.00   sec     0.00 Bytes    0.00 bits/sec    1    1.41 KBytes
[ 12] 7.00-8.00   sec     0.00 Bytes    0.00 bits/sec    0    1.41 KBytes
[ 12] 8.00-9.00   sec     0.00 Bytes    0.00 bits/sec    0    1.41 KBytes
[ 12] 9.00-10.00  sec     0.00 Bytes    0.00 bits/sec    0    1.41 KBytes
-----
[ ID] Interval           Transfer     Bandwidth       Retr
[ 12] 0.00-10.00  sec    76.4 KBytes  62.5 Kbits/sec    5             sender
[ 12] 0.00-10.00  sec     0.00 Bytes    0.00 bits/sec                             receiver

iperf Done.
root@ubuntu:~/cs3611/lab4#

```

Bandwidth is too small and there's no packet received.

This because of the MTU.

From this figure we can get MTU of br1.

```

[sudo] password for dajiaohuang:
ovs-vsctl: cannot create a bridge named br1 because a bridge named br1 already exists
SIOCADDRT: File exists
dajiaohuang@ubuntu:~/cs3611/lab4$ ifconfig -a
br1: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 192.168.56.101 netmask 255.255.255.0 broadcast 192.168.56.255
    inet6 fe80::ccdf:91ff:fe8:814c prefixlen 64 scopeid 0x20<link>
    ether 00:0c:29:5b:f7:51 txqueuelen 1000 (Ethernet)
    RX packets 500 bytes 39932 (39.9 KB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 530 bytes 44064 (44.0 KB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

ens33: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 192.168.186.133 netmask 255.255.255.0 broadcast 192.168.186.255
    inet6 fe80::3799:7062:cd69:6b68 prefixlen 64 scopeid 0x20<link>
    ether 00:0c:29:5b:f7:51 txqueuelen 1000 (Ethernet)
    RX packets 2740 bytes 612248 (612.2 KB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 1303 bytes 141816 (141.8 KB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
    loop txqueuelen 1000 (Local Loopback)
    RX packets 9789 bytes 770301 (770.3 KB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 9789 bytes 770301 (770.3 KB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

ovs-system: flags=4098<BROADCAST,MULTICAST> mtu 1500
    ether c2:8f:7d:bf:f9:52 txqueuelen 1000 (Ethernet)
    RX packets 0 bytes 0 (0.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

s1: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 10.0.0.3 netmask 255.0.0.0 broadcast 10.255.255.255
    inet6 fe80::3064:5aff:feab:2d4d prefixlen 64 scopeid 0x20<link>
    ether 32:64:5a:ab:2d:4d txqueuelen 1000 (Ethernet)
    RX packets 62 bytes 7644 (7.6 KB)
    RX errors 0 dropped 7 overruns 0 frame 0
    TX packets 53 bytes 7928 (7.9 KB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

s2: flags=4098<BROADCAST,MULTICAST> mtu 1500
    ether ba:93:b7:8b:83:49 txqueuelen 1000 (Ethernet)
    RX packets 0 bytes 0 (0.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

s1-eth1: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet6 fe80::9c7c:e8ff:fe28:294f prefixlen 64 scopeid 0x20<link>
    ether 9e:7c:e8:28:29:4f txqueuelen 1000 (Ethernet)
    RX packets 57 bytes 5066 (5.0 KB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 183 bytes 24361 (24.3 KB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

vuln-ovs-4798: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 65536

```

The MTU of br1 is 1500.

Since ping is an application layer command, we know that a datagram length is consist of 20 bytes datagram header, 20 bytes TCP segment header and application message. And we need 50 bytes for VXLAN in this lab. So the MSS of application message is: $1500 - 20 - 20 - 50 = 1410$.

And local mtu is 65536, which is much bigger. So we waste much time packing and unpacking and in congestion.

4. Designating the MTU size of iperf3:

```
"Node: s1" (root)
root@ubuntu:~/cs3611/lab4# iperf -c 10.0.0.4 -M 1140
WARNING: attempt to set TCP maximum segment size to 1140, but got 536
-----
Client connecting to 10.0.0.4, TCP port 5001
TCP window size: 85.3 KByte (default)
-----
[ 11] local 10.0.0.3 port 34050 connected with 10.0.0.4 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 11] 0.0-10.0 sec  2.07 GBytes  1.78 Gbits/sec
root@ubuntu:~/cs3611/lab4#
```

```
Node: s1 (root)
root@ubuntu:~/cs3611/lab4# iperf -c 10.0.0.4 -M 200
WARNING: attempt to set TCP maximum segment size to 200, but got 536
-----
Client connecting to 10.0.0.4, TCP port 5001
TCP window size: 85.3 KByte (default)
-----
[ 11] local 10.0.0.3 port 55256 connected with 10.0.0.4 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 11] 0.0-10.0 sec  330 MBytes  276 Mbits/sec
root@ubuntu:~/cs3611/lab4# iperf -c 10.0.0.4 -M 1100
WARNING: attempt to set TCP maximum segment size to 1100, but got 536
-----
Client connecting to 10.0.0.4, TCP port 5001
TCP window size: 85.3 KByte (default)
-----
[ 11] local 10.0.0.3 port 48464 connected with 10.0.0.4 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 11] 0.0-10.0 sec  1.82 GBytes  1.56 Gbits/sec
root@ubuntu:~/cs3611/lab4# iperf -c 10.0.0.4 -M 1500
WARNING: attempt to set TCP maximum segment size to 1500, but got 536
-----
Client connecting to 10.0.0.4, TCP port 5001
TCP window size: 85.3 KByte (default)
-----
[ 11] local 10.0.0.3 port 55560 connected with 10.0.0.4 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 11] 0.0-10.3 sec  77.8 KBytes  61.6 Kbits/sec
root@ubuntu:~/cs3611/lab4#
```


Bandwidth is improved. Because making the MTU of iperf3 smaller than that of br1 and ens33 (1500) avoids congestion, packing and unpacking.

Designating the MTU of the two ens33 interfaces:

```
dajiaohuang@ubuntu:~$
dajiaohuang@ubuntu:~$ sudo ifconfig ens33 mtu 1400
dajiaohuang@ubuntu:~$ ifconfig ens33
ens33: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1400
    inet 192.168.186.133 netmask 255.255.255.0 broadcast 192.168.186.255
    inet6 fe80::e25e:7541:f848:f492 prefixlen 64 scopeid 0x20<link>
    inet6 fe80::3799:7062:cd69:6b68 prefixlen 64 scopeid 0x20<link>
    inet6 fe80::7c59:9681:168d:947d prefixlen 64 scopeid 0x20<link>
    ether 00:0c:29:5b:f7:51 txqueuelen 1000 (Ethernet)
    RX packets 1640496 bytes 185539193 (185.5 MB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 18915846 bytes 22794290179 (22.7 GB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

dajiaohuang@ubuntu:~$
dajiaohuang@ubuntu:~$ sudo ifconfig ens33 mtu 5000
dajiaohuang@ubuntu:~$ ifconfig ens33
ens33: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 5000
    inet 192.168.186.133 netmask 255.255.255.0 broadcast 192.168.186.255
    inet6 fe80::e25e:7541:f848:f492 prefixlen 64 scopeid 0x20<link>
    inet6 fe80::3799:7062:cd69:6b68 prefixlen 64 scopeid 0x20<link>
    inet6 fe80::7c59:9681:168d:947d prefixlen 64 scopeid 0x20<link>
    ether 00:0c:29:5b:f7:51 txqueuelen 1000 (Ethernet)
    RX packets 1640872 bytes 185563790 (185.5 MB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 18916022 bytes 22794330989 (22.7 GB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

dajiaohuang@ubuntu:~$
```

```
"Node: s1" (root)
root@ubuntu:~/cs3611/lab4# iperf -c 10.0.0.4 -M 1140
WARNING: attempt to set TCP maximum segment size to 1140, but got 536
-----
Client connecting to 10.0.0.4, TCP port 5001
TCP window size: 85.3 KByte (default)
-----
[ 11] local 10.0.0.3 port 34050 connected with 10.0.0.4 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 11] 0.0-10.0 sec  2.07 GBytes  1.78 Gbits/sec
root@ubuntu:~/cs3611/lab4# iperf -c 10.0.0.4
-----
Client connecting to 10.0.0.4, TCP port 5001
TCP window size: 85.3 KByte (default)
-----
[ 11] local 10.0.0.3 port 58722 connected with 10.0.0.4 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 11] 0.0-10.4 sec  77.8 KBytes  61.5 Kbits/sec
root@ubuntu:~/cs3611/lab4#
```

Not improved. Because we can never modify this MTU to make it bigger than the local MTU to avoid packing and unpacking. Also MTU of bridge is not changed, so actually there's no real difference.